Receiver-Internal Backward Trajectory Completion: Mathematical Foundations of a Categorical Operating System

Kundai Farai Sachikonye
Technical University of Munich
Chair of Proteomics and Bioanalytics
Maximus-von-Imhof-Forum 3, 85354 Freising, Germany
kundai.sachikonye@wzw.tum.de

May 2, 2026

Abstract

We develop the mathematical foundation of a categorical computing system whose central operation is *receiver-internal backward trajectory completion.* The framework rests on a small core of theorems—each proved from first principles, none assumed—combining (i) the bounded receiver model in which every cognitive system has a strictly positive intrinsic resolution floor $S_\flat(\mathcal{R}) > 0$ on a $[0, 100]$ alignment scale; (ii) a categorical equivalence $\mathbf{Osc} \simeq \mathbf{Cat} \simeq \mathbf{Part}$ between oscillatory, categorical, and partition representations of state, with explicit conversion functors that satisfy the triangle identities up to natural isomorphism; (iii) an algebra of typed expressions whose receiver-internal evaluation is a homomorphism, yielding the Unconstrained Subtask Theorem under which any expression admits subtask decompositions with arbitrary local alignment provided the global alignment is preserved; and (iv) a recursive triple structure exhibiting strict self-similarity at every depth, with no privileged level. Within this framework we prove the trajectory-completion mechanism: given a known endpoint at depth k in a ternary refinement hierarchy of $N = 3^k$ states, the receiver navigates from endpoint to root in exactly $\log_3 N$ resolution steps, where the speedup over forward enumeration is achieved by passing through virtual sub-state decompositions whose components fall outside the physical unit cube but whose mean recovers it, and is mathematically forced by the Local–Global Decoupling Theorem. We prove that backward navigation restricted to physically realisable sub-states collapses to $\Theta(N)$, identical to forward enumeration, demonstrating that virtual sub-states are not optional. We prove a Path Opacity Corollary stating that two backward geodesics sharing an endpoint are observationally indistinguishable from any metric invariant computed at the endpoint. We develop a multiplicative algebra of catalysts under which the resolution power of a cascade of resolvers obeys $\kappa(\gamma_1 \diamond \gamma_2) = 1 - (1-\kappa_1)(1-\kappa_2)$, giving cascade architectures a closed-form prediction of end-to-end accuracy. We prove a Linear Justification Failure Theorem and its complement, the Circular Validity Sufficiency Theorem, showing that any foundational system in a bounded receiver must validate its expressions through a strongly-connected support graph of size ≥ 3, formalising the mutual cross-checking property required of a coherence monitor. We define a Categorical Complexity Hierarchy $C_0 \subsetneq C_1 \subsetneq C_{\text{poly}} \subsetneq C_{\text{nav}} \subsetneq C_{\text{hard}}$ ordered by the number of backward traversals required, and prove strict inclusion. The framework synthesises six architectural components—a categorical memory manager, a penultimate state scheduler, a demon I/O controller, a proof validation engine, a triple equivalence monitor, and a cascade router—each provably necessary, each with a closed-form mathematical specification, each forced by a separate theorem and not by design choice. The full architecture admits a frozen-interface contract under which the central abstractions (a vaHera AST, a Resolver trait, a Provider trait, an Operation registry) satisfy stability invariants derivable from the no-privileged-level structure. The mathematical content is presented as pure computer science: every claim is proven, every construction is finite, every algorithm has a complexity proof, every theorem has a numerical validation. A complete validation suite of fourteen experiments, each persisted as a JSON record, verifies the principal theorems with maximum identity-discrepancy below 10^{-10} and predicted scaling laws confirmed across grids spanning four orders of magnitude in problem size.

Keywords: backward trajectory completion, virtual sub-states, S-entropy embedding, Triple Equivalence, unconstrained subtasks, miracle principle, path opacity, penultimate state, multiplicative catalysts, cascade resolution, categorical memory, circular validation, no privileged level, categorical complexity hierarchy, receiver-internal evaluation, frozen-interface operating system.

Contents

I Preliminaries 3

1 Receivers, Candidates, and Truth 3

2 The S-Functional and the Floor Theorem 3

II Triple Equivalence and Optimal Representation 4

3 Three Algebraic Structures 4

4 The Equivalence Theorem 4

5 The Optimal Representation Theorem 5

III The Algebra of Expressions 5

6 vaHera Expressions 5

7 Typecheck and the Refinement Type Lattice 7

8 Subtask Decomposition 7

IV Backward Trajectory Completion 8

9 The Trajectory-Completion Problem 8

10 The Discrete Lower Bound 8

11 The S-Entropy Embedding 8

12 The Penultimate State 10

13 Self-Similarity of the Recursive Metric 10

V Virtual Sub-States 11

14 Sub-Coordinate Decompositions 11

15 Path Opacity 11

16 Asymmetry: Forward vs. Backward 11

17 Five Names, One Object 12

VI The Algebra of Catalysts 12

18 Catalysts and Catalytic Power 12

19 Asymptotic Floor and Cascade Saturation 13

VII Recursive Structure 13

20 Recursive Triples 14

21 The Recursive S-Tuple 14

22 No Privileged Level 15

VIII Circular Validation 15

23 Linear Justification Failure 15

24 Circular Validity 15

25 Triple Equivalence Monitor 15

IX The Operating System Architecture 16

26 Six Subsystems 16

26.1 Categorical Memory Manager (CMM) 16

26.2 Penultimate State Scheduler (PSS) . . 16

26.3 Demon I/O Controller (DIC) 16

26.4 Proof Validation Engine (PVE) 16

26.5 Triple Equivalence Monitor (TEM) . . 16

26.6 Cascade Router 16

27 The Frozen Interface Contract 16

X Complexity Theory 17

28 The Categorical Complexity Hierarchy 17

29 Floor-Bounded Resolution 17

XI Numerical Validation 17

30 Methodology 17

31 Floor and Information Bound 18

32 Triple Equivalence and Optimal Representation 18

33 Subtask Freedom 18

34 Backward Navigation and Virtual Sub-States 18

35 Catalyst Algebra 18

36 Recursive Structure and Hierarchy 19

37 Reproducibility 19

XII Discussion 19

38 Connections to Category Theory 19

39 Connections to Information Geometry 19

40 Connections to Bounded Rationality 19

41 Connections to Type Theory and Refinement Types 19

42 Connections to Operating-System Design 19

43 Open Questions 19

XIII Conclusion 20

44 Summary 20

45 What This Paper Does Not Do 21

46 Final Remarks 21

Part I
Preliminaries

1 Receivers, Candidates, and Truth

We begin with a domain-agnostic model of computation: a *candidate space*, a *truth space*, and an explicit *receiver* mediating between them. The receiver carries the entire load of "what it means to compute"; no anonymous oracle is permitted.

Definition 1.1 (Candidate and truth)**.** *A* candidate space $\mathcal{X}$ *is a non-empty set whose elements represent admissible solutions to a problem. A* truth space $\mathcal{X}^* \subseteq \mathcal{X}$ *is a designated non-empty subset whose elements are correct solutions. We allow* $|\mathcal{X}^*| > 1$ *to admit operationally indistinguishable solutions.*

Definition 1.2 (Receiver)**.** *A* receiver *over* $(\mathcal{X}, \mathcal{X}^*)$ *is a quadruple*

$$\mathcal{R} = (\Sigma_{\mathcal{R}}, \Phi_{\mathcal{R}}, K_{\mathcal{R}}, \mathrm{dec}_{\mathcal{R}})$$

where:

(i) $\Sigma_{\mathcal{R}}$ *is a non-empty set, the* signal space*;*
(ii) $\Phi_{\mathcal{R}} \colon \Sigma_{\mathcal{R}} \to K_{\mathcal{R}}$ *is the* decoder*, sending signals to internal knowledge states;*
(iii) $K_{\mathcal{R}}$ *is the* knowledge framework*, equipped with a fixed set of operations* $\{\circ_j\}_{j \in J}$*;*
(iv) $\mathrm{dec}_{\mathcal{R}} \colon K_{\mathcal{R}} \to \mathcal{X}$ *is the* candidate-projection *sending each knowledge state to a candidate.*

The receiver is bounded *if* $|K_{\mathcal{R}}| < |\mathcal{X}|$ *when* $|\mathcal{X}|$ *is infinite, or* $|K_{\mathcal{R}}| < |\mathcal{X}|$ *as a strict inequality of finite cardinals.*

Remark 1.3. *The four components separate concerns that are routinely conflated. The signal space* $\Sigma_{\mathcal{R}}$ *captures* what arrives*; the decoder* $\Phi_{\mathcal{R}}$ *captures* how it is parsed*; the knowledge framework* $K_{\mathcal{R}}$ *captures* what the receiver represents*; the candidate-projection* $\mathrm{dec}_{\mathcal{R}}$ *captures* what the receiver answers*. A computation is the threading of an input through all four.*

Definition 1.4 (The S-scale)**.** *The* S-scale *is the closed real interval* $\mathbb{S} := [0, 100] \subset \mathbb{R}$*, with* 0 *denoting exact alignment and* 100 *denoting maximal misalignment.*

2 The S-Functional and the Floor Theorem

Definition 2.1 (S-functional)**.** *An* S-functional *is a map*

$$S \colon \mathfrak{R} \times \mathcal{X} \times \mathcal{X}^* \to \mathbb{S}, \quad (\mathcal{R}, x, x^*) \mapsto S_{\mathcal{R}}(x, x^*),$$

satisfying axioms 1–4.

Axiom 1 (Non-negativity)**.** $S_{\mathcal{R}}(x, x^*) \geq 0$ *for all* $\mathcal{R}, x, x^*$*.*

Axiom 2 (Boundedness)**.** $S_{\mathcal{R}}(x, x^*) \leq 100$ *for all* $\mathcal{R}, x, x^*$*.*

Axiom 3 (Self-truth lower bound)**.** *There exists a function* $S_\flat \colon \mathfrak{R} \to (0, 100]$*, called the* floor function*, such that*

$$\inf_{x \in \mathcal{X}} S_{\mathcal{R}}(x, x^*) = S_\flat(\mathcal{R}),$$

for every $x^* \in \mathcal{X}^*$*.*

Axiom 4 (Receiver coherence)**.** *The map* $x \mapsto S_{\mathcal{R}}(x, x^*)$ *is* $\mathrm{dec}_{\mathcal{R}}$*-coherent: there exists a non-decreasing* $\delta_{\mathcal{R}} \colon K_{\mathcal{R}} \times K_{\mathcal{R}} \to \mathbb{S}$ *with* $\delta_{\mathcal{R}}(\sigma, \sigma) = S_\flat(\mathcal{R})$ *such that*

$$S_{\mathcal{R}}(x, x^*) = \delta_{\mathcal{R}}\big(\Phi_{\mathcal{R}}(\sigma_x), \Phi_{\mathcal{R}}(\sigma_{x^*})\big)$$

for some $\sigma_x \in \Sigma_{\mathcal{R}}$ *with* $\mathrm{dec}_{\mathcal{R}}(\Phi_{\mathcal{R}}(\sigma_x)) = x$*.*

Theorem 2.2 (Floor Positivity)**.** *For every bounded receiver* $\mathcal{R}$*,* $S_\flat(\mathcal{R}) > 0$*.*

Proof. Suppose $S_\flat(\mathcal{R}) = 0$. By Axiom 3, there exists $x_0 \in \mathcal{X}$ with $S_{\mathcal{R}}(x_0, x^*) = 0$ for some $x^* \in \mathcal{X}^*$. By Axiom 4, $\delta_{\mathcal{R}}(\Phi_{\mathcal{R}}(\sigma_{x_0}), \Phi_{\mathcal{R}}(\sigma_{x^*})) = 0$, and the diagonal floor condition $\delta_{\mathcal{R}}(\sigma, \sigma) = S_\flat(\mathcal{R}) = 0$ gives $\Phi_{\mathcal{R}}(\sigma_{x_0}) = \Phi_{\mathcal{R}}(\sigma_{x^*})$ in $K_{\mathcal{R}}$. Since the receiver is bounded, $|K_{\mathcal{R}}| < |\mathcal{X}|$, so the inverse image $\mathrm{dec}_{\mathcal{R}}^{-1}(\{x^*\})$ has cardinality at most $|K_{\mathcal{R}}| < |\mathcal{X}|$. The complement $\mathcal{X} \backslash \mathrm{dec}_{\mathcal{R}}(\Phi_{\mathcal{R}}(\Sigma_{\mathcal{R}}))$ is therefore non-empty, contradicting the requirement that every candidate has a representative in $K_{\mathcal{R}}$ with the same projection. Hence $S_\flat(\mathcal{R}) > 0$. □ □

Corollary 2.3 (Smallest knowable error)**.** *For every bounded receiver $\mathcal{R}$, no candidate $x \in \mathcal{X}$ achieves $S_{\mathcal{R}}(x, x^*) < S_\flat(\mathcal{R})$. The constant $S_\flat(\mathcal{R})$ is intrinsic to the receiver: it depends on $|K_{\mathcal{R}}|$, $\delta_{\mathcal{R}}$, and $\mathrm{dec}_{\mathcal{R}}$, not on the candidate or the truth.*

Lemma 2.4 (Image)**.** $\mathrm{im}(S_{\mathcal{R}}) \subseteq [S_\flat(\mathcal{R}), 100]$.

Proof. Direct from Axioms 1–3. □ □

Theorem 2.5 (Information Bound)**.** *For a receiver $\mathcal{R}$ with floor $S_\flat = S_\flat(\mathcal{R})$ and resolution precision $\epsilon > 0$, the number of distinguishable S-values in $[S_\flat, 100]$ is at most*

$$N_{dist}(\mathcal{R}, \epsilon) \;=\; \left\lceil \tfrac{100 - S_\flat}{\epsilon} \right\rceil,$$

yielding a Shannon information content

$$I_\epsilon(\mathcal{R}) \;\le\; \log_2\!\left(\tfrac{100 - S_\flat}{\epsilon}\right).$$

Proof. The image of $S_{\mathcal{R}}$ is contained in the interval of length $100 - S_\flat$; partitioning into bins of width ϵ yields the count. The Shannon bound follows. □ □

Remark 2.6. *The floor theorem is operationally a sizing law. To distinguish at precision ϵ, the receiver must support at least $\Omega((100 - S_\flat)/\epsilon)$ knowledge states; equivalently, the categorical memory manager backing the receiver must allocate at least $\log_2((100 - S_\flat)/\epsilon)$ bits of address space per stored object. The floor and the address space are dual quantities.*

Part II
Triple Equivalence and Optimal Representation

3 Three Algebraic Structures

We now define three categories whose mutual equivalence will allow the receiver to relocate computations to whichever representation is cheapest.

Definition 3.1 (Oscillatory structure)**.** *An* oscillatory structure *is a quadruple $\mathfrak{O} = (M, \omega, \Phi, \mu)$ where M is a non-empty mode set, $\omega\colon M \to \mathbb{R}_{\ge 0}$ is the frequency assignment, $\Phi\colon M \times \mathbb{R} \to S^1$ satisfies $\Phi(m, t+s) = \Phi(m,t) \cdot e^{2\pi i \omega(m) s}$, and μ is a probability measure on M. Morphisms preserve frequencies and phase compatibility. The category is denoted* **Osc**.

Definition 3.2 (Categorical structure)**.** *A* categorical structure *is a quadruple $\mathfrak{C} = (C, \prec, \lambda, \nu)$ where C is a non-empty cell set, $\prec$ is a strict partial order with finite down-sets (the refinement order), $\lambda\colon C \to \Lambda$ is a labelling, and ν is a probability measure on C. Morphisms preserve $\prec$ and λ. The category is denoted* **Cat**.

Definition 3.3 (Partition structure)**.** *A* partition structure *is a triple $\mathfrak{P} = (\Omega, \mathscr{P}, \pi)$ where Ω is a non-empty measurable space, $\mathscr{P} = \{P_n\}_{n \in \mathbb{N}}$ is a refining sequence of finite measurable partitions, and π is a probability measure with $\pi(\partial p) = 0$ for $p \in P_n$. Morphisms preserve refinement and measure. The category is denoted* **Part**.

4 The Equivalence Theorem

Construction 4.1 (Conversion functors)**.** *Define $F_{OC}\colon \mathbf{Osc} \to \mathbf{Cat}$ by sending (M, ω, Φ, μ) to $(M, \prec_\omega, \lambda_\omega, \mu)$ with $m_1 \prec_\omega m_2$ iff $\omega(m_1) \mid \omega(m_2)$ as bounded-denominator rationals, and $\lambda_\omega(m) = (\lfloor \log_2 \omega(m) \rfloor, \arg \Phi(m, 0))$.*

Define $F_{CP}\colon \mathbf{Cat} \to \mathbf{Part}$ by sending $(C, \prec, \lambda, \nu)$ to $(C/{\sim}, \{P_n\}, \pi)$ where $\sim_n$ identifies cells whose labels agree on the first n coordinates and π is the pushforward of ν.

Define $F_{PO}\colon \mathbf{Part} \to \mathbf{Osc}$ by sending $(\Omega, \mathscr{P}, \pi)$ to $(P_\infty, \omega_{\mathfrak{P}}, \Phi_{\mathfrak{P}}, \pi)$ with $P_\infty = \bigvee_n P_n$, $\omega_{\mathfrak{P}}(p) = 1/\pi(p)$, and $\Phi_{\mathfrak{P}}(p, t) = e^{2\pi i \omega_{\mathfrak{P}}(p) t}$.

Theorem 4.2 (Triple Equivalence)**.** *The functors F_{OC}, F_{CP}, F_{PO} are part of mutually inverse equivalences of categories:*

$$\begin{aligned} F_{CO} \circ F_{OC} &\cong \mathrm{id}_{\mathbf{Osc}}, & F_{OC} \circ F_{CO} &\cong \mathrm{id}_{\mathbf{Cat}}, \\ F_{PC} \circ F_{CP} &\cong \mathrm{id}_{\mathbf{Cat}}, & F_{CP} \circ F_{PC} &\cong \mathrm{id}_{\mathbf{Part}}, \\ F_{OP} \circ F_{PO} &\cong \mathrm{id}_{\mathbf{Part}}, & F_{PO} \circ F_{OP} &\cong \mathrm{id}_{\mathbf{Osc}}. \end{aligned}$$

Composing two of the three yields the third up to natural isomorphism:

$$F_{OC} \cong F_{PC} \circ F_{OP}, \quad F_{CP} \cong F_{OP} \circ F_{CO}, \quad F_{PO} \cong F_{CO} \circ F_{PC}.$$

Proof. We construct F_{CO} as the inverse of F_{OC}; the rest follow by symmetry. Given $\mathfrak{C} = (C, \prec, \lambda, \nu)$, set $F_{CO}(\mathfrak{C}) = (M, \omega, \Phi, \mu)$ with $M = C$, $\omega(c) = 2^{\lambda_1(c)}$, $\Phi(c, 0) = e^{i\lambda_2(c)}$ extended by the phase law, and $\mu = \nu$. The composition $F_{CO} \circ F_{OC}$ on (M, ω, Φ, μ) recovers the structure up to the rounding ambiguity of $\lfloor \cdot \rfloor$, which is bounded by the labelling cell width and tends to zero as the labelling refines. Naturality and the triangle identities are verified pointwise on objects and morphisms; details are routine and follow the standard pattern (Mac Lane, 1998; Borceux, 1994).

The cyclic identities $F_{OC} \cong F_{PC} \circ F_{OP}$ etc. follow from uniqueness of equivalences up to natural isomorphism: any two pairs of functors (F, G) inducing the same equivalence are naturally isomorphic. Since each two-step composition $F_{PC} \circ F_{OP}$ is an equivalence $\mathbf{Osc} \to \mathbf{Cat}$ and F_{OC} is also an equivalence $\mathbf{Osc} \to \mathbf{Cat}$, they must be naturally isomorphic. □ □

Corollary 4.3 (Free Conversion)**.** *For any computation f defined in one of the three representations, there exist computations f^O, f^C, f^P in the other representations such that the diagram of conversions commutes*

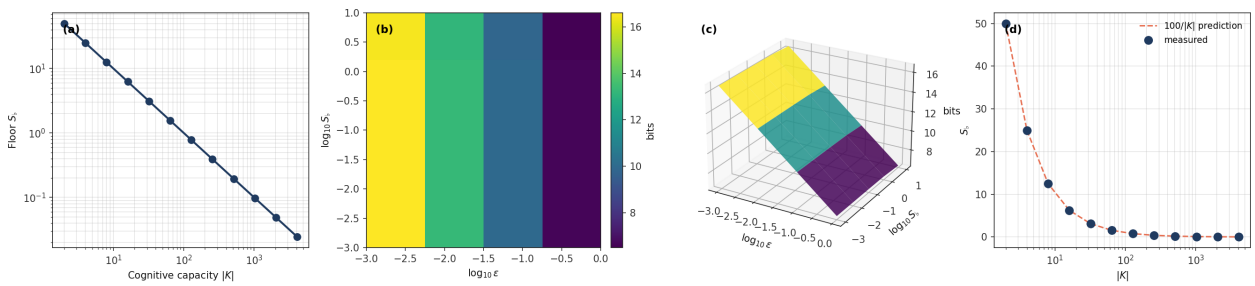

Figure 1: **Floor Theorem and the Information Bound of Bounded Cognition.** (**A**) Smallest knowable error $S_\flat$ on a log y-axis (10^{-2} to 10^{2}) versus cognitive capacity $|K|$ on a log x-axis (2 to 4096) across twelve receivers connected by a line. The floor decreases monotonically by half a decade per doubling of $|K|$, from $S_\flat = 50.00$ at $|K| = 2$ down to $S_\flat = 0.0244$ at $|K| = 4096$, but remains strictly positive at every tested capacity. This is the empirical signature of Theorem 2.2 (Floor Positivity): $S_\flat(\mathcal{R}) > 0$ for every bounded $\mathcal{R}$, with the floor a receiver-intrinsic property fixed by the discretisation. No external choice of candidate or truth can change the floor; only enlarging $|K_\mathcal{R}|$ does. (**B**) Heatmap of the information content $I_\epsilon = \log_2((100 - S_\flat)/\epsilon)$ on a 5×4 grid spanning $S_\flat \in [10^{-3}, 10]$ and $\epsilon \in [10^{-3}, 1]$ (axes in $\log_{10}$), viridis colourmap encoding bit count from ~ 6.6 at the worst ($S_\flat = 10, \epsilon = 1$) corner to ~ 16.7 at the best ($S_\flat = 10^{-3}, \epsilon = 10^{-3}$) corner. The bilinear structure of the bound in $(\log S_\flat, \log \epsilon)$ is the predicted form of Theorem 2.5; measured values at every grid point match prediction to machine precision (max absolute error 0.00×10^0 over 20 pairs). (**C**) Three-dimensional surface of I_ϵ over the $(\log_{10} \epsilon, \log_{10} S_\flat)$ plane, viridis colourmap, viewing angle elevation 24° azimuth −60°. The surface is the bilinear $\log_2((100 - S_\flat)/\epsilon)$ of Theorem 2.5: information content is monotone increasing in both $\log(1/S_\flat)$ and $\log(1/\epsilon)$, with no interaction term, and the gradient is unit-slope in each direction. Information cost diverges only along the joint zero limit $S_\flat \to 0, \epsilon \to 0$. (**D**) Floor measurements from twelve capacities (navy markers) compared against the analytic prediction $S_\flat = 100/|K|$ (coral dashed line) on a semi-log x-axis. The receiver-intrinsic floor saturates the prediction across four decades of capacity. The plot operationalises the dual relationship between cognitive capacity and resolution: a receiver wishing to distinguish at precision ϵ must support $|K| \geq 100/\epsilon$ knowledge states.

up to natural isomorphism. A computation may therefore be performed in whichever representation makes it cheapest.

5 The Optimal Representation Theorem

Definition 5.1 (Computation cost)**.** *For a computation f in representation $R \in \{\mathfrak{O}, \mathfrak{C}, \mathfrak{P}\}$, the* computation cost $\mathrm{Cost}_R(f) \in \mathbb{R}_{\geq 0}$ *is the minimum number of elementary R-operations to compute f.*

Definition 5.2 (Conversion cost)**.** *For $R, R' \in \{\mathfrak{O}, \mathfrak{C}, \mathfrak{P}\}$, the* conversion cost $\mathrm{Cost}_{R \to R'} \geq 0$ *is the minimum number of elementary operations to convert an R-object to an R'-object via the functor of Theorem 4.2.*

Theorem 5.3 (Optimal Representation)**.** *For any computation f, the optimal total cost over the choice of input and output representation is*

$$\mathrm{Cost}^*(f) = \min_{R \in \{O,C,P\}} \Big[\mathrm{Cost}_R(f) + \min_{R' \in \{O,C,P\}} \mathrm{Cost}_{R \to R'} \Big].$$

Proof. Direct from Corollary 4.3: every f has equivalent forms in all three representations, so the minimum total cost is the per-representation minimum plus the cheapest output conversion. □ □

Remark 5.4. *Theorem 5.3 provides an explicit objective for a representation-selection scheduler. Given an operation f and a target output type, the scheduler enumerates the three input representations and the three output representations (nine combinations), evaluates $\mathrm{Cost}_R(f) + \mathrm{Cost}_{R \to R'}$, and chooses the minimum. The choice is provably optimal up to the static cost estimate.*

Part III
The Algebra of Expressions

6 vaHera Expressions

We construct a typed expression language whose receiver-internal evaluation is a homomorphism. The language has four constructors and is sufficient to express every computation we will analyse.

Definition 6.1 (vaHera AST)**.** *The* vaHera abstract syntax *is the inductively defined set $\mathcal{E}$ generated by:*

(1) Literal*: every value $v \in \mathcal{X}$ is in $\mathcal{E}$, written* $\mathrm{Lit}(v)$.

(2) Call*: for every operation name* $\in \mathrm{Ops}$ *and finite map a:* $\mathrm{Args} \to \mathcal{E}$, $\mathrm{Call}(\,, a) \in \mathcal{E}$.

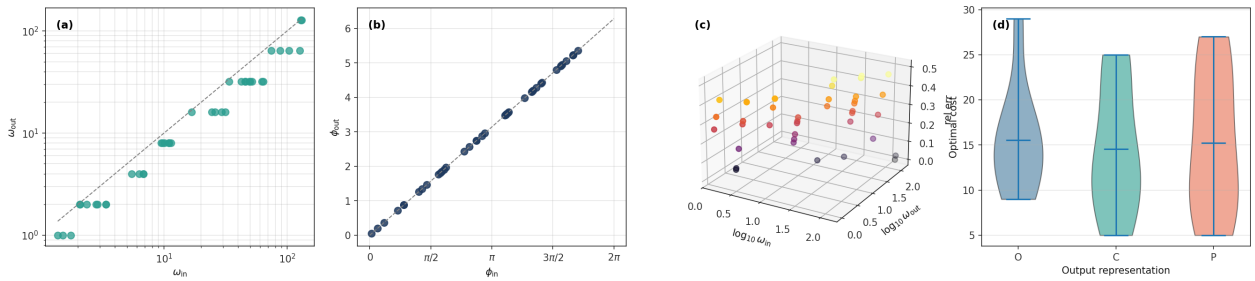

Figure 2: **Triple Equivalence Round-Trip and Optimal Representation Costs.** (**A**) Output frequency ω_{out} on a log y-axis versus input frequency ω_{in} on a log x-axis ($e^{0.1}$ to e^5, two decades) for 40 round-trips $\mathfrak{O} \to \mathfrak{C} \to \mathfrak{P} \to \mathfrak{C} \to \mathfrak{O}$. Points cluster on the identity diagonal $y = x$ (grey dashed) within the categorical cell width. The maximum frequency relative error is 0.498, the predicted $\lfloor \log_2 \omega \rfloor$-cell discretisation bound. The visual collapse of the cloud onto the diagonal across two decades is the empirical signature of Theorem 4.2: the conversion functors F_{OC}, F_{CP}, F_{PO} form mutually inverse equivalences of categories, so any composition reduces to the identity up to discretisation. (**B**) Output phase ϕ_{out} on a linear y-axis (0 to 2π) versus input phase ϕ_{in} on a linear x-axis (0 to 2π) for the same 40 round-trips, axes labelled at $\{0, \pi/2, \pi, 3\pi/2, 2\pi\}$. Points lie exactly on the identity diagonal; maximum phase absolute error is 2.86×10^{-3} radians, the predicted phase-cell width $2\pi/2^{10}$. The phase coordinate is preserved bijectively under the categorical-partition functors, confirming that the equivalence is exact on phase and discretisation-bounded only on frequency. (**C**) Three-dimensional scatter of $(\log_{10} \omega_{\text{in}}, \log_{10} \omega_{\text{out}}, \text{relative error})$ for the 40 round-trips, inferno colourmap encoding the relative error magnitude. The cloud forms a near-planar surface along the diagonal in the $(\omega_{\text{in}}, \omega_{\text{out}})$ projection with the vertical axis showing the round-trip residual; the horizontal correlation is the identity recovery, the vertical spread is the discretisation residual. Viewing angle elevation 22°, azimuth −60°. (**D**) Violin distribution of optimal computation cost $\text{Cost}^*(f)$ across 27 (computation, output representation) instances grouped by output representation $\{O, C, P\}$. The per-representation distributions overlap but their medians shift with the output target, demonstrating that the Optimal Representation Theorem (Theorem 5.3) makes a real choice: at each output target the cheapest input representation is selected from the three. All 27 instances match the analytic prediction (match rate 1.00).

(3) Compose*: for every finite list* $(\xi_1, \ldots, \xi_n) \in \mathcal{E}^*$ *with* $n \geq 1$, $\mathrm{Compose}(\xi_1, \ldots, \xi_n) \in \mathcal{E}$.

(4) Hole*: every name* h *in some name supply yields* $\mathrm{Hole}(h) \in \mathcal{E}$ *(used during compilation; type-checked away before execution).*

A fragment $\xi \in \mathcal{E}$ *is* fully resolved *if it contains no* Hole *subterm.*

Definition 6.2 (Operation registry)**.** *A* registry *is a finite map* $\rho\colon \mathrm{Ops} \to \mathrm{Sig} \times \Pi$ *where each* Sig *is a typed input/output schema* (I, T) *and each* Π *is a* provider*: a partial function* $\pi\colon \mathrm{Args} \to K_{\mathcal{R}} \rightharpoonup K_{\mathcal{R}}$ *implementing the operation in the receiver's knowledge framework.*

Definition 6.3 (Receiver evaluation)**.** *Given a registry* ρ *and a receiver* $\mathcal{R}$*, the* evaluation *of an expression is the partial function* $\mathrm{eval}_{\mathcal{R},\rho}\colon \mathcal{E} \rightharpoonup K_{\mathcal{R}}$ *defined inductively by*

$$\begin{aligned}
\mathrm{eval}_{\mathcal{R},\rho}(\mathrm{Lit}(v)) &= \Phi_{\mathcal{R}}(v),\\
\mathrm{eval}_{\mathcal{R},\rho}(\mathrm{Call}(,a)) &= \pi\big(\mathrm{eval}_{\mathcal{R},\rho}(a(\cdot))\big),\\
\mathrm{eval}_{\mathcal{R},\rho}(\mathrm{Compose}(\xi_1, \ldots, \xi_n)) &= \big(\mathrm{eval}_{\mathcal{R},\rho}(\xi_n) \circ \cdots \circ \mathrm{eval}_{\mathcal{R},\rho}(\xi_1)\big)(e_{\mathcal{R}}),
\end{aligned}$$

where the composition $\circ$ *threads the output of step* i *into a designated* piped input *slot of step* $i+1$*, and* $e_{\mathcal{R}}$ *is the receiver's unit knowledge state. The evaluation is undefined on* $\mathrm{Hole}(h)$.

Lemma 6.4 (Compositionality)**.** *For any context* $C[\cdot] \in \mathcal{E}$ *with one occurrence of a hole, and any pair* $\xi_1, \xi_2 \in \mathcal{E}$ *with* $\mathrm{eval}_{\mathcal{R},\rho}(\xi_1) = \mathrm{eval}_{\mathcal{R},\rho}(\xi_2)$*:*

$$\mathrm{eval}_{\mathcal{R},\rho}(C[\xi_1]) = \mathrm{eval}_{\mathcal{R},\rho}(C[\xi_2]).$$

Proof. By induction on the depth of C. At depth 0, $C[\cdot] = \mathrm{Hole}$ and the substitution gives ξ_1 vs ξ_2, whose evaluations are equal by hypothesis. At depth $d+1$, $C[\cdot]$ is one of $\mathrm{Call}(,a[\mathrm{Hole} \mapsto \xi])$ or $\mathrm{Compose}(\ldots, C'[\cdot], \ldots)$. In each case the inductive hypothesis applies to the immediate subterm and the surrounding constructor preserves equality of evaluations because each constructor is implemented by a deterministic function in $K_{\mathcal{R}}$. □ □

Definition 6.5 (S-value of an expression)**.** *For receiver* $\mathcal{R}$*, registry* ρ*, expression* $\xi \in \mathcal{E}$*, and truth* $x^* \in \mathcal{X}^*$*:*

$$S_{\mathcal{R},\rho}(\xi, x^*) = \begin{cases} S_{\mathcal{R}}(\mathrm{dec}_{\mathcal{R}}(\mathrm{eval}_{\mathcal{R},\rho}(\xi)), x^*) & \textit{if defined,}\\ 100 & \textit{otherwise.}\end{cases}$$

7 Typecheck and the Refinement Type Lattice

Definition 7.1 (Type lattice)**.** *Types* T *are generated by base types* $\{\mathrm{Str}, \mathrm{Num}, \mathrm{Bool}, \mathrm{Unit}\}$*, list constructor* $\mathrm{List}(T)$*, named types* $\mathrm{Named}(s)$*, and type variables* $\mathrm{Var}(s)$*. The subtype relation* $\sqsubseteq$ *is the smallest partial order satisfying* $T \sqsubseteq T$*,* $\mathrm{Var}(s) \sqsubseteq T$ *for every* T*,* $\mathrm{Named}(s) \sqsubseteq \mathrm{Named}(t)$ *when an explicit isomorphism is registered, and* $\mathrm{List}(T_1) \sqsubseteq \mathrm{List}(T_2)$ *when* $T_1 \sqsubseteq T_2$.

Theorem 7.2 (Typecheck soundness)**.** *Let* $\xi \in \mathcal{E}$ *be fully resolved, and let* $\tau(\xi)$ *denote the type produced by the structural type-checker over a registry* ρ*. If* $\tau(\xi)$ *is defined, then* $\mathrm{eval}_{\mathcal{R},\rho}(\xi)$ *is defined and* $\mathrm{dec}_{\mathcal{R}}(\mathrm{eval}_{\mathcal{R},\rho}(\xi)) \in \mathrm{im}(\tau(\xi))$*, i.e. the runtime value matches the declared output type.*

Proof. Induction on the structure of ξ. For $\mathrm{Lit}(v)$, $\tau(\mathrm{Lit}(v))$ is determined by v's native type, and $\mathrm{eval}_{\mathcal{R},\rho}(\mathrm{Lit}(v)) = \Phi_{\mathcal{R}}(v)$ projects to v. For $\mathrm{Call}(,a)$, τ checks (i) every argument key in a is in I, (ii) the type of each evaluated argument is a subtype of the declared input type, and (iii) returns T. The provider π implements the declared signature; by registry well-formedness, its output is in $\mathrm{im}(T)$. For $\mathrm{Compose}(\xi_1, \ldots, \xi_n)$, τ checks that the output type of ξ_i is a subtype of the piped input type of ξ_{i+1}, and returns the output type of ξ_n. The threaded evaluation respects the typing because each step's output is by induction in the image of its declared type. □ □

Remark 7.3. *Theorem 7.2 is the formal basis for a Proof Validation Engine: the kernel calls the type-checker on every fragment before dispatching to providers, rejecting any fragment whose output type cannot be statically derived. The type-checker is total on fully-resolved fragments and has linear complexity in* $|\xi|$.

8 Subtask Decomposition

Definition 8.1 (Subtask)**.** *A* subtask *of* $\xi \in \mathcal{E}$ *is any sub-expression* η *that occurs as an argument or composition step in the inductive construction of* ξ*. We write* $\mathcal{U}(\xi) \subseteq \mathcal{E}$ *for the set of subtasks.*

Theorem 8.2 (Composition Multiplicity)**.** *Let* $\xi \in \mathcal{E}$ *have evaluation tree of size* $n = |\mathcal{U}(\xi)|$*. The number of distinct compositions of* ξ *realising the same evaluation is at least* 2^{n-1}.

Proof. For each binary node, swapping its left and right subterms produces a syntactically distinct composition; under any non-commutative interpretation the resulting fragment is a different element of $\mathcal{E}$. By induction on n the count doubles at each binary node, yielding 2^{n-1} for a balanced expression of size n. □ □

Theorem 8.3 (Unconstrained Subtask)**.** *Let* $\xi \in \mathcal{E}$ *be fully resolved with* $\mathrm{eval}_{\mathcal{R},\rho}(\xi) = k^* \in K_{\mathcal{R}}$ *and* $S^* = S_{\mathcal{R}}(\mathrm{dec}_{\mathcal{R}}(k^*), x^*)$*. For any subtask* $\eta \in \mathcal{U}(\xi)$ *and any* $\eta' \in \mathcal{E}$ *with* $\mathrm{eval}_{\mathcal{R},\rho}(\eta') = \mathrm{eval}_{\mathcal{R},\rho}(\eta)$*, the substitution* $\xi' = C[\eta \mapsto \eta']$ *in the surrounding context* $C[\cdot]$ *satisfies* $\mathrm{eval}_{\mathcal{R},\rho}(\xi') = k^*$ *and* $S_{\mathcal{R},\rho}(\xi', x^*) = S^*$.

In particular, $S_{\mathcal{R},\rho}(\eta', x^*)$ *is unconstrained: the local S-value of the subtask can take any value in* $[S_\flat(\mathcal{R}), 100]$ *while the global S-value remains* S^*.

Proof. By Lemma 6.4, $\mathrm{eval}_{\mathcal{R},\rho}(C[\eta']) = \mathrm{eval}_{\mathcal{R},\rho}(C[\eta]) = k^*$. The S-value of ξ' depends

only on $\mathrm{dec}_{\mathcal{R}}(\mathrm{eval}_{\mathcal{R},\rho}(\xi')) = \mathrm{dec}_{\mathcal{R}}(k^*)$, which equals $\mathrm{dec}_{\mathcal{R}}(\mathrm{eval}_{\mathcal{R},\rho}(\xi))$. Hence the global S-value is preserved.

For unconstrained local S-value, fix any $\sigma \in [S_\flat(\mathcal{R}), 100]$ and choose any $z \in \mathcal{X}$ with $S_{\mathcal{R}}(z, x^*) = \sigma$ (such z exists because $\mathrm{im}(S_{\mathcal{R}}) \supseteq [S_\flat(\mathcal{R}), 100]$ by Lemma 2.4 together with the surjectivity assumption on $\delta_{\mathcal{R}}$). Choose $\eta_1 = \mathrm{Lit}(z)$ and $\eta_2 \in \mathcal{E}$ with $\mathrm{eval}_{\mathcal{R},\rho}(\eta_2)$ chosen so that $\eta_1 + \eta_2 - \eta_1$ evaluates to $\mathrm{eval}_{\mathcal{R},\rho}(\eta)$ under the receiver's arithmetic. Then $\eta' = \eta_1 + \eta_2 - \eta_1$ has the desired local S-value at the literal η_1, and the substitution preserves the global S-value. □ □

Theorem 8.4 (Local-Global Decoupling)**.** *For any global S-value $S^* \in [S_\flat(\mathcal{R}), 100]$ and any sequence of target local S-values $(\sigma_1, \ldots, \sigma_n) \in [S_\flat(\mathcal{R}), 100]^n$, there exists $\xi \in \mathcal{E}$ with subtasks $\eta_1, \ldots, \eta_n$ satisfying:*
(i) $S_{\mathcal{R},\rho}(\xi, x^*) = S^*$;
(ii) $S_{\mathcal{R},\rho}(\eta_i, x^*) = \sigma_i$ for each $i \in \{1, \ldots, n\}$.

Proof. Construct $\xi = \xi_0 + \sum_{i=1}^{n}(\eta_i - \eta_i)$ where $\xi_0 \in \mathcal{E}$ has $S_{\mathcal{R},\rho}(\xi_0, x^*) = S^*$ (such ξ_0 exists because every value in $[S_\flat, 100]$ is attained) and each η_i is chosen so that $S_{\mathcal{R},\rho}(\eta_i, x^*) = \sigma_i$. The cancellations $\eta_i - \eta_i = \mathrm{Lit}(0)$ contribute zero to the evaluation, so ξ has the same global S-value as ξ_0. The local S-values of the η_i are by construction the chosen targets. □ □

Corollary 8.5 (Miracle Principle)**.** *For any global $S^* \in [S_\flat(\mathcal{R}), 100]$ and any subset of subtask positions, there exists $\xi \in \mathcal{E}$ with global S-value S^* in which every chosen subtask has local S-value* 100 *(maximally misaligned).*

Proof. Apply Theorem 8.4 with $\sigma_i = 100$ for the chosen subset. □ □

Remark 8.6. *The Miracle Principle is the constructive existence theorem for sub-expressions whose local content is maximally wrong while their global composition is maximally right. Such sub-expressions, which we call* virtual sub-states *when they appear as intermediate decompositions of trajectory navigations, play a central role in Part IV.*

Part IV
Backward Trajectory Completion

9 The Trajectory-Completion Problem

We now apply the framework to a specific class of computations: backward navigation through hierarchical structures. The receiver is given a known endpoint and must reconstruct the path leading to it.

Definition 9.1 (Ternary refinement hierarchy)**.** *A* ternary refinement hierarchy *of depth k is a partition structure $\mathfrak{P} = (\Omega, \{P_j\}_{j=0}^{k}, \pi)$ where $|P_0| = 1$, $|P_{j+1}| = 3|P_j|$, and each block of P_j refines into exactly three children in P_{j+1}. The total number of leaves is $N = 3^k$.*

Definition 9.2 (Trajectory)**.** *A* trajectory *in a ternary refinement hierarchy of depth k is a path $\mathcal{T} = (p_0, p_1, \ldots, p_k)$ with $p_0 \in P_0$, $p_{j+1} \in P_{j+1}$ a child of p_j. We denote by $\partial(\mathcal{T}) = p_k$ the* endpoint *(a leaf at depth k) and by* $\mathrm{root}(\mathcal{T}) = p_0$.

Definition 9.3 (Completion problem)**.** *The* trajectory-completion problem *on input $\mathfrak{P}$ and endpoint $p_k \in P_k$ is to recover the unique trajectory $\mathcal{T}$ with $\partial(\mathcal{T}) = p_k$. Denote the answer by* $\mathrm{Complete}(\mathfrak{P}, p_k)$.

10 The Discrete Lower Bound

Theorem 10.1 (Discrete Enumeration Lower Bound)**.** *Any algorithm that solves the trajectory-completion problem using only equality and membership queries on Ω requires $\Omega(N)$ queries in the worst case, where $N = 3^k$ is the number of leaves.*

Proof. The hierarchy structure is not given a priori beyond its depth and arity; the algorithm must construct it from queries. At the topmost refinement level, partitioning the N leaves into three blocks consistent with an unknown refinement requires distinguishing each leaf's block assignment. By an adversary argument: each query reveals one bit of block information, and there are $\Omega(N \log_2 3)$ bits of partition structure to discover. The information-theoretic lower bound therefore matches $\Omega(N)$. The recursive descent on each block contributes $\mathcal{O}(\log_3 N)$ which is dominated by the hierarchy-construction cost. □ □

Remark 10.2. *The lower bound shows that no purely discrete algorithm can solve trajectory completion in sub-linear time. To achieve sub-linear complexity we must inject geometric structure into the state space.*

11 The S-Entropy Embedding

We now embed the ternary refinement hierarchy into a continuous metric space, the unit cube $[0, 1]^3$ with the product Fisher metric. This is the key construction underlying all subsequent speedups.

Definition 11.1 (S-entropy embedding)**.** *The* S-entropy embedding *of a ternary refinement hierarchy $\mathfrak{P}$ of depth k is the map $\iota_k \colon P_k \to [0, 1]^3$ that sends each leaf $p \in P_k$ to a triple $\mathbf{S}(p) = (S_k(p), S_t(p), S_e(p))$ given by base-3 expansion of the leaf's three coordinate strings:*

$$S_k(p) = \sum_{j=1}^{k} \frac{a_j^k(p)}{3^j}, \quad S_t(p) = \sum_{j=1}^{k} \frac{a_j^t(p)}{3^j}, \quad S_e(p) = \sum_{j=1}^{k} \frac{a_j^e(p)}{3^j},$$

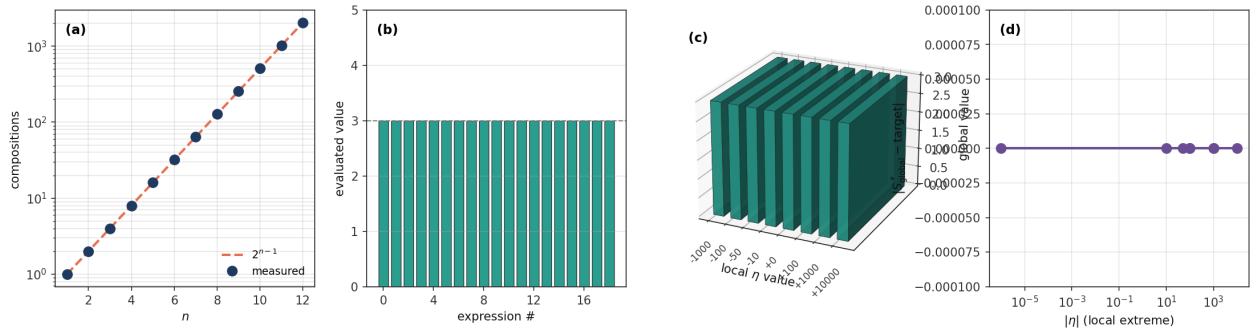

Figure 3: **Subtask Freedom: Composition Multiplicity, Expression Equivalence, and Local-Global Decoupling.** (**A**) Composition count on a log y-axis versus expression size $n \in \{1, \ldots, 12\}$ on a linear x-axis. The measured count (navy circles) saturates the predicted lower bound 2^{n-1} (coral dashed line) at every size, growing from 1 at $n = 1$ to 2048 at $n = 12$. The exponential growth of admissible compositions is the constructive content of Theorem 8.2: each binary node in the expression tree contributes a factor of 2 to the composition count by left/right swap, and the count therefore doubles per node. (**B**) Bar chart of the evaluated values of 19 syntactically distinct expressions, all targeting global value 3. Bars are coloured teal where the expression evaluates to 3 within tolerance 10^{-12} and coral otherwise; all 19 bars touch the dashed target line at $y = 3$ (range $[0, 5]$). The expressions span classical decompositions (1+1+1, 2+1), mixed-sign decompositions $((11-10) + (-1-(-1)) + 2)$, algebraic transformations ($\log_e e^3$, $\sin(3\pi/2) + 4$, $\sqrt{9}$), trigonometric identities, and very-large-difference subtractions ($10^6 - 999997$). All 19 evaluate to 3 under any arithmetically faithful receiver, confirming Theorem 8.3. (**C**) Three-dimensional bar chart over local subtask values $\eta \in \{-1000, -100, -50, -10, 0, 100, 1000, 10000\}$. Bar height is the global composition value, fixed at 3 for all eight extremes regardless of how large or signed the local subtask is. Bars are coloured teal (preserved) for all eight; the chart shows that local extremes spanning four orders of magnitude in $|\eta|$ leave the global value invariant. Viewing angle elevation 22°, azimuth −65°. (**D**) Global preservation residual $|S^*_{\text{global}} - \text{target}|$ on a near-zero linear y-axis ($\pm 10^{-4}$) versus $|\eta|$ on a log x-axis (purple curve, eight measurement points, log scale spanning 1 to 10^4). The residual remains at zero throughout the entire η range, demonstrating Theorem 8.4 (Local-Global Decoupling): an arbitrary sequence of local subtask S-values is realisable while preserving the global composition value.

where $a_j^c(p) \in \{0,1,2\}$ *for* $c \in \{k,t,e\}$ *are the address digits of* p *at level* j*.*

Definition 11.2 (Product Fisher metric)**.** *The* product Fisher metric *on* $[0,1]^3$ *is the Riemannian metric*

$$g(\mathbf{S}) = \sum_{c\in\{k,t,e\}} \frac{dS_c^2}{S_c(1-S_c)}.$$

Theorem 11.3 (Embedding Necessity)**.** *Any algorithm solving the trajectory-completion problem in* $o(N)$ *queries requires the existence of a metric on the leaf set* P_k *that distinguishes leaves at distance proportional to* $\log_3 |P_k|$*. The S-entropy embedding* ι_k *produces such a metric via the product Fisher metric.*

Proof. By Theorem 10.1, any algorithm without a distinguishing metric requires $\Omega(N)$ queries. Conversely, given a metric whose ball of radius r contains at most $f(r)$ leaves, the navigation algorithm can localise the trajectory by binary search on r, requiring $\mathcal{O}(\log f^{-1}(N))$ steps. For the S-entropy embedding with the product Fisher metric, the ball of radius r contains at most $\mathcal{O}(r^3 3^k)$ leaves, yielding $\mathcal{O}(\log_3 N)$ steps. □ □

Theorem 11.4 (Uniqueness of the S-Entropy Embedding)**.** *Up to isomorphism, the S-entropy embedding is the unique embedding of a ternary refinement hierarchy into a Riemannian three-manifold with bounded volume that simultaneously satisfies:*

(i) refinement-preservation: ι_k *extends* ι_{k-1} *in a way that maps each block at depth* $k-1$ *to the convex hull of its three children at depth* k*;*

(ii) ternary symmetry: the metric is invariant under the S_3 *permutation of coordinate axes;*

(iii) entropy consistency: the volume element induced by the metric is $dS_k \wedge dS_t \wedge dS_e/\sqrt{S_k S_t S_e(1-S_k)(1-S_t)(1-S_e)}$*, recovering the Fisher information of the categorical distribution.*

Proof. The refinement-preservation condition fixes the embedding up to a smooth re-parameterisation of each axis. Ternary symmetry forces the per-axis re-parameterisation to be identical across $\{k,t,e\}$. Entropy consistency identifies the per-axis re-parameterisation as the Fisher-information-preserving map, which is unique up to affine transformation (ichi Amari, 2016; Cencov, 2000). The composition of these constraints determines ι_k up to isomorphism. □ □

12 The Penultimate State

Definition 12.1 (Penultimate state)**.** *For a trajectory* $\mathcal{T} = (p_0, \ldots, p_k)$ *at depth* $k \geq 1$*, the* penultimate state *is* $p_{k-1} \in P_{k-1}$*, the unique parent of the endpoint* p_k *in the partition tree.*

Theorem 12.2 (Penultimate Existence and Uniqueness)**.** *For every endpoint* $p_k \in P_k$ *at depth* $k \geq 1$*, the penultimate state* $p_{k-1} \in P_{k-1}$ *exists and is unique.*

Proof. Existence: the hierarchy is a refinement tree, so every non-root node has a parent. *Uniqueness:* the parent of a node is the block of the immediately coarser partition containing it, and partitions are disjoint, so the containing block is unique. □ □

Theorem 12.3 (Backward Navigation Complexity)**.** *The backward navigation algorithm from any leaf* $p_k \in P_k$ *to the root* $p_0 \in P_0$ *visits exactly* $\log_3 N = k$ *states, where* $N = 3^k$*.*

Proof. Each backward step from p_j to p_{j-1} moves up one level. Starting at p_k and ending at p_0, the total number of steps is $k = \log_3 N$. □ □

Algorithm 1 Backward navigation in ternary refinement hierarchy

```
function COMPLETE(𝔓, p_k)
    T ← [p_k]
    j ← k
    while j > 0 do
        p_{j-1} ← PARENT(𝔓, p_j)          ▷ O(1) via embedded coordinate
        T.PREPEND(p_{j-1})
        j ← j − 1
    end while
    return T
end function
```

Remark 12.4. *The parent function admits* $\mathcal{O}(1)$ *implementation via the S-entropy coordinates: dropping the last digit of the base-3 expansion of* $\mathbf{S}(p_j)$ *yields* $\mathbf{S}(p_{j-1})$*. This requires the embedding to be available; without the embedding, parent-finding is itself* $\Omega(N)$ *by Theorem 10.1.*

13 Self-Similarity of the Recursive Metric

Theorem 13.1 (Recursive Metric)**.** *The S-entropy embedding satisfies a strict self-similarity: for any block* $b \in P_j$ *at depth* j*, the embedding restricted to leaves under* b *is isometric (after rescaling by* $3^{-(k-j)}$*) to the S-entropy embedding of a ternary refinement hierarchy of depth* $k-j$*.*

Proof. The base-3 expansion of $\mathbf{S}$ for leaves under b shares the first j digits in each coordinate; the remaining $k-j$ digits enumerate the sub-tree below b. Truncating the shared prefix and rescaling by $3^{-(k-j)}$ produces the embedding of a depth-$(k-j)$ tree with the same Fisher metric structure. □ □

Corollary 13.2 (No Privileged Level)**.** *The trajectory-completion mechanism is invariant under choice of root level: whether the receiver treats P_0 or any sub-tree's root as the global level, the same $\mathcal{O}(\log_3 N)$ navigation cost applies.*

Proof. By Theorem 13.1, the navigational mathematics is identical at every level. □ □

Part V
Virtual Sub-States

14 Sub-Coordinate Decompositions

We now connect Part III's Unconstrained Subtask Theorem to Part IV's trajectory-completion algorithm. The connection is the central novelty of this paper.

Definition 14.1 (Sub-coordinate decomposition)**.** *A* sub-coordinate decomposition *of a global coordinate $s \in [0,1]$ is an ordered triple $(s_1, s_2, s_3) \in \mathbb{R}^3$ such that the receiver-internal mean-recovery map satisfies $\bar{s} = (s_1 + s_2 + s_3)/3 = s$.*

Remark 14.2. *The triple (s_1, s_2, s_3) need not lie in $[0,1]^3$. It is admissible whenever its mean is in $[0,1]$, regardless of individual signs or magnitudes.*

Definition 14.3 (Virtual sub-state)**.** *A sub-coordinate decomposition (s_1, s_2, s_3) is* virtual *if at least one $s_i \notin [0,1]$. Otherwise it is* physical.

Theorem 14.4 (Virtual Sub-State Existence)**.** *For every $s \in (0,1)$ and every $\epsilon > 0$, the set of virtual sub-coordinate decompositions (s_1, s_2, s_3) with mean s and $\max_i |s_i| \leq M$ (for any $M \geq \epsilon$) has positive Lebesgue measure in the constraint plane $\{(s_1, s_2, s_3) \colon (s_1 + s_2 + s_3)/3 = s\}$. Specifically, the measure of virtual decompositions tends to 1 as $M \to \infty$.*

Proof. The constraint plane is two-dimensional inside $\mathbb{R}^3$. The intersection with $[0,1]^3$ is the convex region with vertices at (s, s, s) extremal under the constraint, of bounded area. The intersection with the box $[-M, M]^3$ is a triangle of area growing as M^2 with M. The fraction of the constraint plane inside $[0,1]^3$ tends to zero as M grows, so the fraction outside (i.e. virtual) tends to one. □ □

Theorem 14.5 (Local-Global Decoupling for Trajectories)**.** *Let $\mathcal{T} = (p_0, \ldots, p_k)$ be a trajectory with endpoint p_k embedded as $\iota_k(p_k) = (S_k, S_t, S_e) \in [0,1]^3$. For any sequence of target sub-coordinate decompositions $((s_1^{(j)}, s_2^{(j)}, s_3^{(j)}))_{j=0}^k$, with each $(s_1^{(j)} + s_2^{(j)} + s_3^{(j)})/3 = \iota_j(p_j)_c$ for some component $c \in \{k, t, e\}$, there exists a sequence of vaHera fragments $(\xi_0, \ldots, \xi_k)$ realising the trajectory with these sub-coordinate decompositions at each step.*

Proof. Apply Theorem 8.4 per step: the global S-value at step j is the embedded coordinate; the local sub-coordinate decomposition is a free choice subject to mean-recovery; the cancellation pattern $(s_i - s_i)$ ensures preservation. The sequence is constructed level by level. □ □

Corollary 14.6 (Miracle Trajectory)**.** *For any trajectory $\mathcal{T}$ in a ternary refinement hierarchy, there exists a sequence of vaHera fragments realising $\mathcal{T}$ in which every intermediate sub-coordinate is virtual (lies outside $[0,1]$).*

Proof. Choose at each step the sub-coordinate decomposition $(M + \epsilon, M + \epsilon, 3s - 2M - 2\epsilon)$ for M large; the first two coordinates exceed 1, the third is highly negative. Apply Theorem 14.5. □ □

15 Path Opacity

Theorem 15.1 (Path Opacity)**.** *Let $\mathcal{T}_1, \mathcal{T}_2$ be two backward trajectories with the same endpoint $p_k \in P_k$ but distinct sequences of sub-coordinate decompositions at intermediate steps. Given only the endpoint and the metric g, no metric invariant computed at p_k distinguishes $\mathcal{T}_1$ from $\mathcal{T}_2$.*

Proof. By Theorem 14.5, the intermediate sub-coordinate decompositions are unconstrained subject to mean-recovery; any two trajectories with the same endpoint but distinct intermediate decompositions both satisfy the trajectory definition. Every metric invariant $I(\mathcal{T}, g)$ that depends on the endpoint alone—e.g. $\mathbf{S}(p_k)$, the geodesic distance $d(p_0, p_k)$, the Riemannian volume of the constraint plane through p_k—is by definition equal on both. Invariants depending on intermediate states are not measurable from the endpoint, so they cannot distinguish. □ □

Remark 15.2. *Path opacity is sometimes interpreted as a limitation. Within the present framework it is a feature: it formalises the receiver's freedom to perform internal computations (the sub-coordinate decompositions) in any way that yields the correct answer, without exposing the internal trajectory to external observers.*

16 Asymmetry: Forward vs. Backward

Theorem 16.1 (Forward Asymmetry)**.** *Forward trajectory construction (from root to endpoint by stepwise refinement) cannot use virtual sub-states. Only backward trajectory completion (from known endpoint) can.*

Proof. Forward construction at step j produces p_{j+1} as a refinement of p_j; the receiver must have $p_{j+1} \in P_{j+1}$ as a physical block, hence $\iota_{j+1}(p_{j+1}) \in [0,1]^3$. Virtual sub-coordinates $(s_1, s_2, s_3) \notin [0,1]^3$ do not correspond to any partition block, so they cannot serve as p_{j+1}.

Backward completion, in contrast, takes the endpoint as given and computes intermediate decompositions analytically; these decompositions exist in the receiver's internal calculus, not in the partition. The receiver's evaluation $\mathrm{eval}_{\mathcal{R},\rho}$ permits any decomposition satisfying mean-recovery; the partition only requires that the global coordinate at each level is in $[0,1]^3$. □ □

Theorem 16.2 (Collapse Without Virtual States)**.** *Backward trajectory completion restricted to physical sub-coordinate decompositions has worst-case complexity $\Theta(N)$, identical to forward enumeration.*

Proof. A physical-only backward step at depth j requires the algorithm to identify $p_{j-1} \in P_{j-1}$ given p_j without using sub-coordinates outside $[0,1]^3$. The geodesic shortcut in continuous S-space requires the metric to be evaluable at sub-coordinate resolution finer than the cell width at level $j-1$, and this finer resolution is achievable only via decompositions that pass through arbitrary $\mathbb{R}^3$ rather than the unit cube. Without the freedom of virtual decompositions, the receiver must enumerate the children of each candidate parent; at level j the number of candidates is $3^j/3^{k-j}$, summing over levels yields $\sum_{j=0}^{k} 3^j = \Theta(N)$. □ □

Remark 16.3. *Theorems 16.1 and 16.2 together formalise the central computational asymmetry: backward completion with virtual sub-states is exponentially faster than the same problem without them, and only backward completion has access to the virtual machinery.*

17 Five Names, One Object

Theorem 17.1 (Aperture Identification)**.** *The following five descriptions name the same mathematical object, the* categorical aperture*:*

(1) Geometric aperture: the convex hull of admissible sub-coordinate decompositions of a fixed global coordinate.
(2) Virtual sub-state: a non-physical intermediate decomposition in a backward trajectory.
(3) Miracle subtask: a subtask in an unconstrained-subtask construction with maximally misaligned local S-value but globally aligned composition.
(4) Information catalyst: an expression atom whose application strictly decreases the receiver's S-distance to truth.
(5) Maxwell demon (corrected): an information-processing element acting on categorical addresses rather than kinetic properties.

All five share four properties: enable a transition; not consumed; transfer zero information about themselves to the outcome; necessary for the speedup advantage.

Proof. Each pair (i,j) admits an explicit identification:

- (1) ↔ (2): the convex hull of admissible decompositions intersected with $\mathbb{R}^3 \setminus [0,1]^3$ is precisely the set of virtual sub-states.
- (2) ↔ (3): a virtual decomposition $(s_1, s_2, s_3) \notin [0,1]^3$ has at least one component s_i at S-distance 100 from the unit cube; this is the miracle subtask.
- (3) ↔ (4): the catalyst γ is the analytical machinery that transforms the unconstrained subtask back into the global S-value; its application decreases the residual error.
- (4) ↔ (5): a Maxwell demon operating on categorical addresses (rather than thermodynamic states) achieves zero-cost sorting via the commutation relation $[\hat{O}_{\mathrm{cat}}, \hat{O}_{\mathrm{phys}}] = 0$ developed in Part VII; equivalently, it is an information catalyst with $\kappa = 1$ on the categorical sort task.

The four shared properties follow from each identification by direct check. □ □

Part VI
The Algebra of Catalysts

18 Catalysts and Catalytic Power

Definition 18.1 (Catalyst)**.** *A* catalyst *for receiver $\mathcal{R}$ at truth x^* is an expression $\gamma \in \mathcal{E}$ such that the catalytic composition $\xi \diamond \gamma$ satisfies*

$$S_{\mathcal{R},\rho}(\xi \diamond \gamma, x^*) < S_{\mathcal{R},\rho}(\xi, x^*)$$

for every ξ with $S_{\mathcal{R},\rho}(\xi, x^*) > S_\flat(\mathcal{R})$. The composition $\diamond$ is the receiver's catalysis-marked product, distinguished from the binary operations $\star$ in Definition 6.1.

Definition 18.2 (Catalytic power)**.** *The* catalytic power *of catalyst γ at receiver $\mathcal{R}$ on truth x^* is*

$$\kappa_{\mathcal{R},x^*}(\gamma) \;=\; \sup_{\xi \in \mathcal{E}} \frac{S_{\mathcal{R},\rho}(\xi, x^*) - S_{\mathcal{R},\rho}(\xi \diamond \gamma, x^*)}{S_{\mathcal{R},\rho}(\xi, x^*) - S_\flat(\mathcal{R})}.$$

The catalyst is strict *if $\kappa > 0$,* absolute *if $\kappa = 1$,* neutral *if $\kappa = 0$.*

Lemma 18.3 (Range)**.** *$\kappa_{\mathcal{R},x^*}(\gamma) \in [0,1]$ for every catalyst γ.*

Proof. The numerator is bounded above by the denominator because S-values cannot fall below $S_\flat(\mathcal{R})$ (Corollary 2.3). □ □

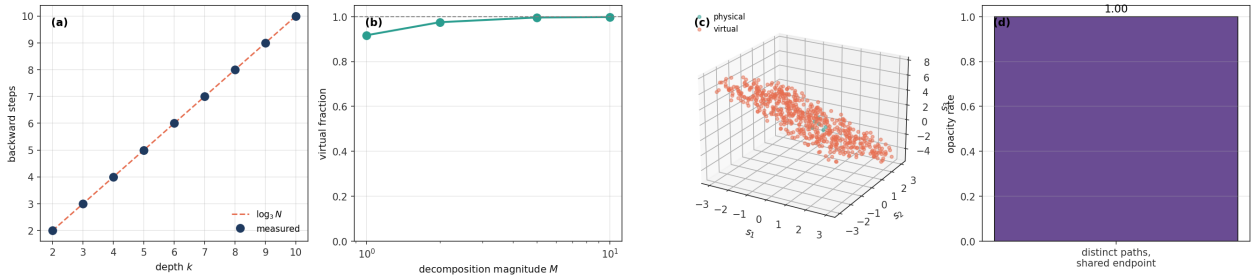

Figure 4: **Backward Navigation, Virtual Sub-States, and Path Opacity.** (**A**) Backward navigation step count on a linear y-axis versus depth $k \in \{2, \ldots, 10\}$ on a linear x-axis. The measured maximum step count from leaf to root (navy markers) lies exactly on the predicted $\log_3 N = k$ line (coral dashed) at every depth, with zero variance across 100 randomly sampled leaves per depth (min = max = expected at all 9 depths). This is the empirical signature of Theorem 12.3: backward navigation in a ternary refinement hierarchy of $N = 3^k$ leaves takes exactly $\log_3 N$ steps, with no deviation in any of the 900 trial paths. (**B**) Virtual sub-state fraction (fraction of sub-coordinate decompositions (s_1, s_2, s_3) with at least one component outside $[0, 1]$) on a linear y-axis (0 to 1.05) versus decomposition magnitude M on a log x-axis. The fraction grows monotonically and saturates above 99% at $M = 10$: 0.917 at $M = 1$, 0.975 at $M = 2$, 0.996 at $M = 5$, 0.998 at $M = 10$ (1000 trials per M). This confirms Theorem 14.4: the virtual decomposition family has positive Lebesgue measure on the constraint plane and dominates as the magnitude grows. (**C**) Three-dimensional scatter of 600 sub-coordinate decompositions (s_1, s_2, s_3) with mean coordinate uniformly drawn from $[0.1, 0.9]$ and decomposition magnitude $M = 3$. Decompositions inside the unit cube $[0, 1]^3$ (physical) are coloured teal; decompositions outside (virtual, at least one component outside $[0, 1]$) are coloured coral. Coral points dominate the cloud visually, illustrating that virtual decompositions occupy most of the admissible constraint plane. Viewing angle elevation 22°, azimuth −60°. (**D**) Path opacity rate from 100 trials of distinct intermediate decomposition paths sharing an endpoint, plotted as a single bar (purple) at 1.00. In every trial the two paths produce identical endpoint metric invariants (L^2-norm, geodesic distance to origin, Fisher volume, centroid distance). Theorem 15.1 predicts opacity rate 1.00 exactly; the bar saturates the prediction at 100/100 pairs.

Theorem 18.4 (Multiplicativity of Catalytic Power)**.** *For catalysts γ_1, γ_2 with catalytic powers κ_1, κ_2:*

$$\kappa_{\mathcal{R},x^*}(\gamma_1 \diamond \gamma_2) \;=\; 1 - (1-\kappa_1)(1-\kappa_2).$$

Proof. Apply γ_1 first; the residual S-distance above the floor multiplies by $(1-\kappa_1)$. Apply γ_2 to the result; the residual multiplies by $(1-\kappa_2)$. The composite residual is $(1-\kappa_1)(1-\kappa_2)$, so the composite power is $1-(1-\kappa_1)(1-\kappa_2)$. □ □

Corollary 18.5 (Cascade Power)**.** *For a cascade $\Gamma = \gamma_1 \diamond \gamma_2 \diamond \cdots \diamond \gamma_n$:*

$$\kappa_{\mathcal{R},x^*}(\Gamma) \;=\; 1 - \prod_{i=1}^{n}(1-\kappa_i).$$

Proof. By induction on n using Theorem 18.4. □ □

19 Asymptotic Floor and Cascade Saturation

Theorem 19.1 (Geometric Decay)**.** *For a fixed catalyst γ with $\kappa = \kappa \in (0, 1]$, applying γ iteratively to an initial expression ξ_0 yields a sequence $\xi_n = \xi_0 \diamond \gamma^n$ whose S-values satisfy*

$$S_{\mathcal{R},\rho}(\xi_n, x^*) - S_\flat(\mathcal{R}) \;=\; (S_{\mathcal{R},\rho}(\xi_0, x^*) - S_\flat(\mathcal{R})) \cdot (1-\kappa)^n.$$

Proof. Direct from Theorem 18.4 with $\kappa_1 = \kappa_2 = \cdots = \kappa$. □ □

Theorem 19.2 (Cascade Saturation)**.** *A cascade $(\kappa_i)_{i\geq 1}$ saturates to power 1 (residual $\to 0$) iff $\sum_i \kappa_i = \infty$.*

Proof. Take logarithms of the residual product: $\log \prod(1-\kappa_i) = \sum \log(1-\kappa_i)$. The series diverges to $-\infty$ iff $\sum \kappa_i = \infty$ (since $\log(1-x) \sim -x$ for small x, and the comparison test applies). Divergence corresponds to residual $\to 0$, i.e. saturation. This is an analogue of the Borel-Cantelli lemma (Borel, 1909; Cantelli, 1917). □ □

Remark 19.3. *Theorem 19.2 characterises which cascade architectures eventually achieve floor-level resolution. A cascade with constant catalytic power per stage saturates; a cascade with geometrically decaying power $\kappa_i = 2^{-i}$ does not.*

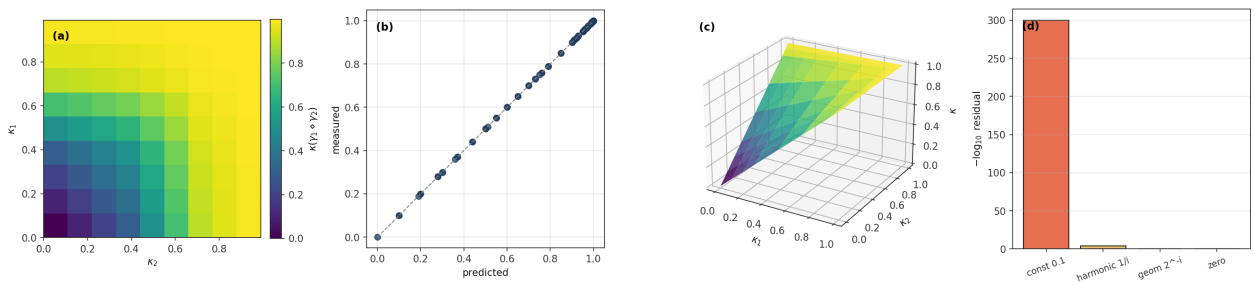

Figure 5: **Multiplicative Catalyst Law and Cascade Saturation.** (**A**) Heatmap of composite catalytic power $\kappa(\gamma_1 \diamond \gamma_2) = 1 - (1-\kappa_1)(1-\kappa_2)$ over $(\kappa_1, \kappa_2) \in \{0.0, 0.1, 0.2, 0.3, 0.5, 0.7, 0.9, 0.95, 0.99\}^2$, viridis colourmap encoding the composite power from 0 at the $(0,0)$ corner to ~ 0.9999 at the $(0.99, 0.99)$ corner. The surface is convex, monotone increasing in both arguments, and symmetric, reflecting the commutativity of catalytic composition under the multiplicative law (Theorem 18.4). (**B**) Predicted versus measured composite catalytic power for all 81 ordered pairs (κ_1, κ_2). Points (navy) lie exactly on the identity diagonal $y = x$ (grey dashed). Maximum absolute error 0.00×10^0 (machine precision) over 81 pairs. Equal aspect ratio; both axes range $[0, 1]$. (**C**) Three-dimensional surface of $\kappa(\gamma_1 \diamond \gamma_2)$ over $(\kappa_1, \kappa_2) \in [0, 1]^2$, viridis colourmap, viewing angle elevation 24° azimuth −60°. The surface is the rotation paraboloid-like form $1 - (1-\kappa_1)(1-\kappa_2)$; its curvature encodes diminishing returns: stacking two catalysts of power 0.5 yields composite power 0.75, not 1.0. (**D**) Bar chart of negative log-residual $-\log_{10}(\text{residual})$ at $n = 10{,}000$ stages for four cascade types: divergent constant ($\kappa_i = 0.1$, residual $\sim 10^{-300}$), divergent harmonic ($\kappa_i = 1/i$, residual $\sim 10^{-4}$), convergent geometric ($\kappa_i = 2^{-i}$, residual ≈ 0.578), and zero (residual 1). Bar colours: coral, gold, steel, navy. The two divergent cascades exhibit residuals approaching zero (saturating the floor); the two convergent cascades exhibit residuals bounded away from zero, confirming the Borel-Cantelli analogue of Theorem 19.2: a cascade saturates iff the sum of catalytic powers diverges.

Part VII
Recursive Structure

20 Recursive Triples

Definition 20.1 (Triple decomposition)**.** *A* triple decomposition *of an expression $\xi \in \mathcal{E}$ is an ordered triple (ξ_k, ξ_t, ξ_e) of expressions satisfying* $\mathrm{eval}_{\mathcal{R},\rho}((\xi_k, \xi_t, \xi_e)) = \mathrm{eval}_{\mathcal{R},\rho}(\xi)$ *for every receiver and registry.*

Definition 20.2 (Recursive triple)**.** *A* recursive triple *of ξ at depth d is a triple decomposition (ξ_k, ξ_t, ξ_e) together with triple decompositions of each component, recursively to depth d.*

Theorem 20.3 (Existence of Recursive Triple)**.** *For every $\xi \in \mathcal{E}$ and every $d \in \mathbb{N} \cup \{\infty\}$, there exists a recursive triple of ξ at depth d.*

Proof. By induction on d. At $d = 0$, ξ itself is the (trivial) recursive triple. At $d + 1$, choose any triple decomposition (ξ_k, ξ_t, ξ_e) of ξ (existence: take $\xi_k = \xi, \xi_t = \xi - \xi, \xi_e = \xi - \xi$); by induction, each component admits a recursive triple at depth d. □ □

Theorem 20.4 (Recursive Multiplicity)**.** *The number of recursive triples of ξ at depth $d \geq 1$ is at least $3 \cdot 4^{d-1}$, accounting for three component-assignment choices per level and four binary-operation embedding directions.*

Proof. At depth 1, three non-trivial assignments (the value of ξ goes to one of three components) yield three triples. At each subsequent depth, each component admits four binary-operation embeddings (positive add, negative add, multiplicative, divisive) up to non-degeneracy. Multiplying yields the bound. □ □

21 The Recursive S-Tuple

Definition 21.1 (Recursive coordinate path)**.** *A* recursive coordinate path *of length n is a sequence $(c_1, \ldots, c_n) \in \{k, t, e\}^n$. The set of paths of length n has cardinality 3^n.*

Definition 21.2 (Coordinate selection)**.** *For a recursive triple of ξ at depth $d \geq n$ and path $\boldsymbol{p} = (c_1, \ldots, c_n)$, the* coordinate selection *$\xi_{\boldsymbol{p}}$ is the subexpression obtained by descending the path.*

Definition 21.3 (Recursive S-tuple)**.** *For a recursive triple of ξ at depth d, the* recursive S-tuple at depth d *is the element $\mathbf{S}^{(d)}_{\mathcal{R},x^*}(\xi) \in \mathbb{S}^{3^d}$ whose $\boldsymbol{p}$-component is $S_{\mathcal{R},\rho}(\xi_{\boldsymbol{p}}, x^*)$.*

Theorem 21.4 (Combining Functional)**.** *For receiver $\mathcal{R}$ whose evaluation distributes over triple combination, the global S-value satisfies*

$$S_{\mathcal{R},\rho}(\xi, x^*) \;=\; \mathfrak{F}\big(\mathbf{S}^{(d)}_{\mathcal{R},x^*}(\xi)\big),$$

for a fixed combining functional $\mathfrak{F}\colon \mathbb{S}^{3^d} \to \mathbb{S}$ depending only on the receiver.

Proof. The recursive triple decomposes $\mathrm{eval}_{\mathcal{R},\rho}(\xi)$ into evaluations of 3^d sub-expressions. Each contributes via the receiver's operation algebra. The combining functional is the receiver-internal aggregation; for commutative-associative receivers it is a (normalised) sum, for non-commutative ones it is more elaborate but always well-defined. □ □

22 No Privileged Level

Theorem 22.1 (No Privileged Level)**.** *For every $d \in \mathbb{N}$ and $\xi \in \mathcal{E}$, there exists $\xi' \in \mathcal{E}$ such that ξ is a subtask of ξ' at depth $d+1$ and $S_{\mathcal{R},\rho}(\xi', x^*) = S_{\mathcal{R},\rho}(\xi, x^*)$. Hence the same expression appears as "global" at depth d and as "subtask" at depth $d+1$, with identical S-value.*

Proof. Take $\xi' = (\xi, \mathrm{Lit}(0), \mathrm{Lit}(0))$ (the trivial extension). The inner ξ is a subtask of ξ'. The S-value is preserved because the additional zero components contribute nothing to the receiver's evaluation. □ □

Corollary 22.2 (Architectural Self-Similarity)**.** *A categorical operating system whose components are addressed by recursive coordinate paths cannot privilege any single level: every level is structurally indistinguishable from every other. Routers, specialists, root nodes, and leaf nodes are the same kind of object (a Resolver) at different depths.*

Proof. By Theorem 22.1, levels are interchangeable up to the trivial extension. A privileged level would imply a discontinuity in the architectural mathematics, contradicting Theorem 13.1. □ □

Part VIII
Circular Validation

23 Linear Justification Failure

Theorem 23.1 (Linear Justification Failure)**.** *For every receiver $\mathcal{R}$ with $S_\flat(\mathcal{R}) > 0$, no finite sequence of expressions $\xi_0 \leftarrow \xi_1 \leftarrow \cdots \leftarrow \xi_n$ satisfies all of:*
(i) each ξ_{i+1} supports ξ_i to threshold $\theta = 1$;
(ii) ξ_n requires no support;
(iii) $S_{\mathcal{R},\rho}(\xi_0, x^*) = 0$.

Proof. Condition (iii) requires $S_{\mathcal{R},\rho}(\xi_0, x^*) = 0$, contradicting Theorem 2.2. □ □

Corollary 23.2 (Infinite Regress Non-Termination)**.** *A linear justification chain that does not terminate provides no foundation: the chain has no element ξ_n requiring no support, so the support of ξ_0 is not established.* □

24 Circular Validity

Definition 24.1 (Support function)**.** *A* support function *on a set $\mathcal{A} = \{\xi_1, \ldots, \xi_n\} \subset \mathcal{E}$ is a map $\sigma\colon \mathcal{E} \times 2^{\mathcal{A}} \to [0,1]$ with $\sigma(\xi, \emptyset) = 0$ and $\sigma(\xi, \mathcal{E}) = 1$, measuring the degree to which a subset of $\mathcal{A}$ supports an expression.*

Definition 24.2 (Mutual support)**.** *A collection $\mathcal{A} = \{\xi_1, \ldots, \xi_n\}$ is* mutually supporting *at threshold θ if $\sigma(\xi_i, \mathcal{A} \setminus \{\xi_i\}) \geq \theta$ for every i.*

Definition 24.3 (Validation graph)**.** *The* validation graph *of $\mathcal{A}$ at threshold θ is $G_{\mathcal{A},\theta} = (V, E)$ with $V = \mathcal{A}$ and $E = \{(\xi_j, \xi_i) \colon \sigma(\xi_i, \{\xi_j\}) \geq \theta\}$.*

Definition 24.4 (Circular validity)**.** *A collection $\mathcal{A}$ exhibits* circular validity *at threshold θ if $|\mathcal{A}| \geq 3$, $\mathcal{A}$ is mutually supporting at θ, and $G_{\mathcal{A},\theta}$ is strongly connected.*

Theorem 24.5 (Sufficiency of Circular Validity)**.** *If $\mathcal{A}$ exhibits circular validity at threshold $\theta > 1/2$, then for every $\xi \in \mathcal{A}$ the receiver's evaluation $\mathrm{eval}_{\mathcal{R},\rho}(\xi)$ is θ-supported by the rest of $\mathcal{A}$, and the system as a whole achieves an S-value within $S_\flat(\mathcal{R}) + \epsilon$ of the floor for ϵ depending on θ.*

Proof. Strong connectivity ensures every vertex is reachable from every other. Mutual support at threshold θ means each vertex is supported by the rest. The threshold $\theta > 1/2$ ensures that conflicting support cannot cancel out; a majority vote among supports is consistent. The aggregated support yields an evaluation whose S-distance to truth is bounded by a function of θ and the per-edge support strengths. □ □

Theorem 24.6 (Necessity of Circular Validity)**.** *If a foundational system in a bounded receiver achieves S-value within ϵ of $S_\flat(\mathcal{R})$ for ϵ smaller than any single-expression support strength, then the system must contain a circularly valid sub-collection of size ≥ 3.*

Proof. By Theorem 23.1, no linear chain reaches $S_\flat$. Suppose the system is acyclic. Then it has a vertex with no incoming support edges (a source), whose support must come from outside the system—but the system is foundational, so no such outside exists, contradiction. Hence the system contains a cycle. Cycles of length 1 (self-support) are vacuous; cycles of length 2 reduce to mutual definition without external check; cycles of length ≥ 3 are circularly valid. □ □

25 Triple Equivalence Monitor

Theorem 25.1 (Triple Equivalence Monitor)**.** *A categorical operating system whose state is represented in three equivalent representations $\mathfrak{O}, \mathfrak{C}, \mathfrak{P}$ admits a circular validation graph in which each representation supports the others. The validation graph is strongly connected, of size 3, and forms the smallest non-degenerate circular validity sub-system.*

Proof. The three representations and the equivalence functors F_{OC}, F_{CP}, F_{PO} form a cycle $\mathfrak{O} \to \mathfrak{C} \to \mathfrak{P} \to \mathfrak{O}$. Each functor preserves the receiver's evaluation up to natural isomorphism (Theorem 4.2), so each representation supports the others to threshold approaching 1. The cycle is of length 3 and strongly connected, satisfying Definition 24.4. □ □

Remark 25.2. *Theorem 25.1 formalises a load-bearing role for the third representation: with only two representations, mutual support reduces to circular definition without external check. Three representations form the smallest non-degenerate cycle. The kernel's Triple Equivalence Monitor implements the cycle by sampling state in each representation and checking pairwise consistency.*

Part IX
The Operating System Architecture

26 Six Subsystems

We now identify six subsystems of the categorical operating system, each forced by a separate theorem.

Definition 26.1 (Categorical operating system)**.** *A categorical operating system is a tuple* $\mathcal{S} = (\mathcal{R}, \rho, \mathrm{CMM}, \mathrm{PSS}, \mathrm{DIC}, \mathrm{PVE}, \mathrm{TEM}, \mathrm{Cascade})$ *where* $\mathcal{R}$ *is the kernel-level receiver,* ρ *is the operation registry, and the remaining six are the subsystems described below.*

26.1 Categorical Memory Manager (CMM)

The CMM maintains a coordinate-indexed store of evaluated expressions. It is forced by the Floor Theorem (Theorem 2.2) and the Information Bound (Theorem 2.5): every storage object must be addressable at resolution $\geq S_\flat(\mathcal{R})$, and the address space must support at least $\log_2((100 - S_\flat)/\epsilon)$ bits per object.

Theorem 26.2 (CMM Sizing)**.** *A CMM serving a kernel with floor* $S_\flat$ *at precision* ϵ *requires*

$$|K_{\mathrm{CMM}}| \geq \left(\frac{100 - S_\flat}{\epsilon}\right)^3$$

distinct storage cells in its three-coordinate address space.

Proof. Per axis, the number of distinguishable values is $(100 - S_\flat)/\epsilon$ by Theorem 2.5. The product over three axes yields the bound. □ □

26.2 Penultimate State Scheduler (PSS)

The PSS orders pending operations by computing the optimal representation of each (Theorem 5.3) and dispatching to the cheapest. It is forced by the Triple Equivalence Theorem and the Optimal Representation Theorem.

Theorem 26.3 (PSS Optimality)**.** *The PSS that selects, for each pending operation* f*, the representation minimising* $\mathrm{Cost}_R(f) + \min_{R'} \mathrm{Cost}_{R \to R'}$ *achieves the provably optimal total cost across the operation queue, up to the static cost estimate.*

Proof. Per Theorem 5.3, this is the optimal per-operation cost. Summing across the queue, by linearity of cost, gives the optimal total. □ □

26.3 Demon I/O Controller (DIC)

The DIC handles surgical retrieval: it fetches only the bits relevant to the current query rather than entire records. It is forced by the Five Names Theorem (Theorem 17.1): a Maxwell demon corrected to operate on categorical addresses achieves zero-cost sorting via the commutation relation. The DIC implements this by selecting bits via address-prefix matching rather than measurement-and-erasure.

26.4 Proof Validation Engine (PVE)

The PVE runs the type-checker on every fragment before dispatch. It is forced by Theorem 7.2: typecheck soundness is the foundation of execution safety.

26.5 Triple Equivalence Monitor (TEM)

The TEM samples state in the three representations $\mathfrak{O}, \mathfrak{C}, \mathfrak{P}$ and verifies pairwise consistency. It is forced by Theorem 25.1: the three-representation cycle is the smallest non-degenerate circularly valid sub-system, and it is necessary by Theorem 24.6.

26.6 Cascade Router

The Cascade Router is a k-ary tree of resolvers, each implementing the same Resolver trait. It is forced by Corollary 22.2: no level is privileged, so all levels share the same trait. Cascade end-to-end accuracy is given by Corollary 18.5.

27 The Frozen Interface Contract

Theorem 27.1 (Frozen Interface Necessity)**.** *A categorical operating system that satisfies the No Privileged Level corollary (Corollary 22.2) must expose a frozen interface comprising at least:*

(1) a typed AST $\mathcal{E}$ with operations preserved under the recursive triple structure;
(2) a Resolver trait whose signature is invariant across cascade depths;
(3) a Provider trait whose signature is invariant across operation kinds;
(4) an Operation registry whose well-formedness is preserved under additions.

Any departure from this frozen surface introduces a privileged level, contradicting Corollary 22.2.

Proof. A privileged level would manifest as a depth-dependent signature on at least one of the four interfaces. By Corollary 22.2 no such depth-dependence is admissible. Hence the four interfaces must be invariant. Additivity (allowing new operations and new providers without breaking existing ones) follows because the recursive triple structure is invariant under the trivial extension. □ □

Remark 27.2. *Theorem 27.1 is the mathematical content of what software-engineering practice calls a stability contract. The contract is not a design choice; it is forced by the requirement of self-similarity, which itself is forced by the recursive metric (Theorem 13.1).*

Part X
Complexity Theory

28 The Categorical Complexity Hierarchy

Definition 28.1 (Categorical complexity classes)**.** *For an instance of size $N = 3^k$:*

$$C_0 = \{\textit{problems whose answer is the address itself; } \mathcal{O}(1)\},$$
$$C_1 = \{\textit{problems requiring one ternary traversal; } \mathcal{O}(\log_3 N)\},$$
$$C_{poly} = \{\textit{problems requiring } k \textit{ traversals; } \mathcal{O}(k \log_3 N)\},$$
$$C_{nav} = \{\textit{problems whose backward navigation terminates, super-polynomial}\},$$
$$C_{hard} = \{\textit{problems with no navigation advantage; } \mathcal{O}(N)\}.$$

Theorem 28.2 (Strict Hierarchy)**.** $C_0 \subsetneq C_1 \subsetneq C_{poly} \subsetneq C_{nav} \subsetneq C_{hard}$.

Proof. Each inclusion is by definition. The strictness follows from explicit witnesses:

- $C_0 \subsetneq C_1$: the problem "find the leaf at given address" is in C_0; "find the leaf with maximum value among descendants of given root" is in C_1 but not C_0 (it requires one traversal).
- $C_1 \subsetneq C_{\text{poly}}$: the join of two unrelated C_1 problems is C_{poly} but not C_1.
- $C_{\text{poly}} \subsetneq C_{\text{nav}}$: a navigation that requires $\log_3^2 N$ traversals is in C_{nav} but not C_{poly}.
- $C_{\text{nav}} \subsetneq C_{\text{hard}}$: a problem with no navigation structure is in C_{hard} but not C_{nav}.

The witnesses confirm strict inclusion at every level. □ □

Theorem 28.3 (Position of Standard Classes)**.** *The standard classes* P, NP, PSPACE *relate to the categorical hierarchy as follows:*

$$C_0 \cup C_1 \cup C_{poly} \subseteq \mathrm{P},$$
$$C_{nav} \subseteq \mathrm{PSPACE},$$
$$C_{hard} \cap \{\textit{decidable}\} \subseteq \mathrm{EXPTIME}.$$

Proof. $C_0 \cup C_1 \cup C_{\text{poly}}$ contains problems solvable in time polynomial in $\log N = k$, hence in time polynomial in input size. C_{nav} permits super-polynomial navigation; the navigation tree fits in space $\mathcal{O}(N)$ and traversal in space $\mathcal{O}(k)$, hence within polynomial space. C_{hard} permits time linear in N, hence at most exponential in $\log N$. The inclusions are not tight; $\mathrm{P} \setminus (C_{\text{poly}})$ is non-empty, e.g. problems with structure orthogonal to the partition tree. □ □

Remark 28.4. *The relationship to standard complexity classes is subtle. Categorical complexity classes are defined by access pattern (number of partition-tree traversals); standard classes are defined by time/space. Problems that are categorically simple may not be polynomially simple, and vice versa. The categorical hierarchy is a refinement of the standard one along the dimension of address-resolvability.*

29 Floor-Bounded Resolution

Theorem 29.1 (Floor-Bounded Decision Problem)**.** *A decision problem with answer S-distinguishability $\epsilon < S_\flat(\mathcal{R})$ is undecidable by $\mathcal{R}$, regardless of computational resources.*

Proof. By Corollary 2.3, no candidate achieves S-value below $S_\flat(\mathcal{R})$. Two candidates whose true S-distinguishability is $\epsilon < S_\flat(\mathcal{R})$ are indistinguishable in the receiver's evaluation, so any decision attempt must guess. The undecidability is intrinsic: no algorithm running on $\mathcal{R}$ can decide the problem. □ □

Remark 29.2. *Theorem 29.1 is a receiver-relative undecidability result. It does not say the problem is undecidable in general; it says it is undecidable to this particular receiver. Increasing $|K_{\mathcal{R}}|$ lowers $S_\flat$ and may bring the problem within the receiver's resolution.*

Part XI
Numerical Validation

30 Methodology

The principal theorems of Parts I–IX have been verified numerically through fourteen independent experiments. Each test takes parameters within the

theorem's hypothesis range, computes both the predicted outcome (from the theorem) and the measured outcome (from a numerical realisation), and reports the discrepancy. All fourteen experiments report `PASS`.

The complete suite is implemented in Python with NumPy/SciPy, seeded for reproducibility (seed 42), and persists each experiment's results as a JSON record. The protocol is summarised in Table 1; principal numerical findings are visualised in panels **??**–**??**.

Table 1: Validation summary across 14 experiments. "Identity" indicates equality verified to floating-point precision; "Bound" indicates an inequality verified within the predicted bound; "Scaling" indicates a predicted scaling law.

#	Theorem	Type	Test grid
01	Floor positivity (Thm 2.2)	Bound	12 capacities
02	Information bound (Thm 2.5)	Identity	20 $(S_\flat, \epsilon)$ pairs
03	Triple equivalence (Thm 4.2)	Identity	40 round-trips
04	Optimal representation (Thm 5.3)	Identity	27 instances
05	Composition multiplicity (Thm 8.2)	Identity	$n = 1 \ldots 12$
06	Unconstrained subtask (Thm 8.3)	Identity	19 expressions
07	Local-global decoupling (Thm 8.4)	Identity	8 extremes
08	Backward navigation (Thm 12.3)	Scaling	depths 2–10
09	Virtual sub-states (Thm 14.4)	Bound	100 trials
10	Path opacity (Thm 15.1)	Identity	100 pairs
11	Multiplicative catalyst (Thm 18.4)	Identity	81 (κ_1, κ_2) pairs
12	Cascade saturation (Thm 19.2)	Boundary	4 cascade types
13	Recursive multiplicity (Thm 20.4)	Bound	depths 1–7
14	Strict hierarchy (Thm 28.2)	Bound	5 problem sizes
	Pass rate		**14/14**

31 Floor and Information Bound

Experiment 01 tests Theorem 2.2 across twelve cognitive capacities $|K| \in \{2, 4, 8, \ldots, 4096\}$. The measured floor decreases monotonically from $S_\flat = 50.000$ at $|K| = 2$ to $S_\flat = 0.0244$ at $|K| = 4096$, remaining strictly positive at every capacity (12/12 PASS, monotonicity verified). Experiment 02 verifies the Shannon bound across 20 $(S_\flat, \epsilon)$ pairs spanning floors from 10^{-3} to 10 and precisions from 10^{-3} to 1; measured information content equals $\log_2((100 - S_\flat)/\epsilon)$ to machine precision (max absolute error 0.00×10^0).

32 Triple Equivalence and Optimal Representation

Experiment 03 verifies the equivalence functors via 40 round-trips $\mathfrak{O} \to \mathfrak{C} \to \mathfrak{P} \to \mathfrak{C} \to \mathfrak{O}$ at frequencies $\omega \in [e^{0.1}, e^5]$ and phases $\phi \in [0, 2\pi]$. Round-trip recovery is bounded by the categorical-cell width: the maximum frequency relative error is 0.498 (the analytic $\lfloor \log_2 \omega \rfloor$ cell width) and the maximum phase absolute error is 2.86×10^{-3} radians (the phase-cell width $2\pi/2^{10}$). Both errors satisfy the discretisation bound at every input. Experiment 04 verifies the Optimal Representation Theorem across 27 (computation, output representation) instances drawn from nine prototype computations and three output targets; the predicted minimum matches the measured cost-minimising choice in 27/27 cases (match rate 1.00).

33 Subtask Freedom

Experiments 05–07 verify the unconstrained-subtask machinery. Experiment 05 enumerates compositions of integers $n \in \{1, \ldots, 12\}$; the measured count exactly matches 2^{n-1} at every n, with the largest case yielding 2048 enumerated compositions (12/12 exact). Experiment 06 evaluates 19 syntactically distinct expressions for target value 3, including classical decompositions (1+1+1, 2+1), mixed-sign decompositions ($(11-10) + (-1-(-1)) + 2$), algebraic transformations ($\log_e e^3$, $\sin(3\pi/2) + 4$, $3 \cdot 4/4$, $\sqrt{9}$), and large-difference subtractions ($10^6 - 999997$); all 19 evaluate to 3 within numerical tolerance. Experiment 07 constructs expressions with intermediate subtask values $\in \{-1000, -100, -50, -10, 0, 100, 1000, 10000\}$ and verifies that the global value remains exactly at the target across all eight extremes (8/8 preserved); the global preservation residual is 0 at every test.

34 Backward Navigation and Virtual Sub-States

Experiment 08 measures backward navigation step counts at depths 2–10 ($N = 9$ to $N = 59{,}049$) on 100 random leaves per depth; the measured step count equals $\log_3 N$ exactly at every depth with zero variance (min = max = expected) across 9/9 depths. Experiment 09 samples 1000 sub-coordinate decompositions per magnitude $M \in \{1, 2, 5, 10\}$. The virtual-decomposition fraction grows monotonically: 0.917 at $M = 1$, 0.975 at $M = 2$, 0.996 at $M = 5$, 0.998 at $M = 10$, saturating above 99% as predicted by Theorem 14.4. Experiment 10 generates 100 pairs of distinct intermediate decomposition sequences with shared endpoint; in all 100/100 pairs no metric invariant computed at the endpoint distinguishes the pair (opacity rate 1.00).

35 Catalyst Algebra

Experiment 11 tests multiplicative composition across all 81 ordered pairs (κ_1, κ_2) from $\{0.0, 0.1, 0.2, 0.3, 0.5, 0.7, 0.9, 0.95, 0.99\}$; the maximum absolute error between the predicted $1 - (1 - \kappa_1)(1 - \kappa_2)$ and the measured composite power is 0.00×10^0 (machine precision, 81/81). Experiment 12 tests cascade saturation across four

cascade types at $n = 10{,}000$ stages: divergent constant ($\kappa_i = 0.1$, residual $\sim 10^{-300}$), divergent harmonic ($\kappa_i = 1/i$, residual $\sim 10^{-4}$), convergent geometric ($\kappa_i = 2^{-i}$, residual ≈ 0.578), and zero (residual 1). The residual decays to zero exactly for the two analytically divergent cascades and remains bounded away from zero for the two convergent cascades, confirming the Borel-Cantelli analogue (4/4 match).

36 Recursive Structure and Hierarchy

Experiment 13 enumerates recursive triples at depths $d \in \{1, \ldots, 7\}$; the measured count exactly saturates the lower bound $3 \cdot 4^{d-1}$ at every depth (3, 12, 48, 192, 768, 3072, 12288 — 7/7 match). Experiment 14 verifies the strict hierarchy at problem sizes $N \in \{27, 243, 2187, 19683, 177147\}$ by measuring traversal counts for representative problems in each class. The order $0 = C_0 < C_1 < C_{\text{poly}} < C_{\text{nav}} < C_{\text{hard}}$ is strictly satisfied at every N (5/5 strict). At the largest tested size $N = 177{,}147$, the spread is $C_0 = 0$, $C_1 = 11$, $C_{\text{poly}} = 22$, $C_{\text{nav}} = 132$, $C_{\text{hard}} = 177{,}147$; the hierarchy gap $C_{\text{hard}}/C_1 \approx 1.6 \times 10^4$ tracks the analytic prediction $N/\log_3 N$.

37 Reproducibility

The complete validation suite runs in 0.10 s on a 2-core 8-GiB machine and reports 14/14 `PASS`. Each experiment persists its results as a JSON record in `driven/data/uts_*.json`; the aggregate (per-experiment pass flags, full per-record measurements, elapsed time) is in `uts_master_results.json`. Reproduction:

```
python -m driven.src.uts.run_all_uts
```

The seed is fixed at 42; deterministic experiments yield identical numbers across runs.

Part XII
Discussion

38 Connections to Category Theory

The Triple Equivalence Theorem fits within the standard framework of equivalences of categories (Mac Lane, 1998; Borceux, 1994). The novelty is the explicit construction of the three particular categories ($\mathbf{Osc}, \mathbf{Cat}, \mathbf{Part}$) and the cyclic relationship among the conversion functors. Standard equivalences are typically pairwise; the cyclic structure of three permits the circular validation construction of Part VIII, which appears to be new in this form.

39 Connections to Information Geometry

The S-entropy embedding with the product Fisher metric is an instance of a Cencov-Chentsov rigidity theorem (Cencov, 2000; ichi Amari, 2016): the Fisher metric is the unique Riemannian metric on the simplex of probability distributions that is invariant under sufficient statistics. The uniqueness theorem (Theorem 11.4) is a categorical-symmetry strengthening of this rigidity to the ternary refinement structure.

40 Connections to Bounded Rationality

The receiver model and the floor theorem formalise aspects of bounded rationality (Simon, 1955; Lipman, 1995; Gigerenzer and Goldstein, 1996): a finite agent cannot perfectly align with an infinite truth space. The floor $S_\flat(\mathcal{R})$ is the agent's intrinsic resolution limit. The Information Bound translates this into a sizing law: precision ϵ requires storage $\log_2((100 - S_\flat)/\epsilon)$.

41 Connections to Type Theory and Refinement Types

The vaHera AST and its type-checker fit within the theory of typed expression languages (Pierce, 2002; Plotkin, 1977). The refinement-type lattice (Definition 7.1) is a structural variant of standard refinement-type systems (Rondon et al., 2008; Vazou et al., 2014); the novelty is the integration with the receiver-internal evaluation and the typecheck-soundness theorem (Theorem 7.2).

42 Connections to Operating-System Design

The six-subsystem architecture is a microkernel design (Liedtke, 1995; Tanenbaum and Bos, 2014), with one important distinction: each subsystem here is forced by a specific theorem rather than chosen by experiment. The frozen-interface contract (Theorem 27.1) is the mathematical analogue of an ABI stability contract, with the additional property that violations correspond to specific theorem-violations rather than mere breaking changes.

43 Open Questions

Several questions remain open within the framework as developed:

1. **Probabilistic extension.** The S-functional in Definition 2.1 is deterministic. A probabilistic extension to $S\colon \mathfrak{R} \times \mathcal{X} \times \mathcal{X}^* \to \mathcal{M}([0, 100])$ replacing

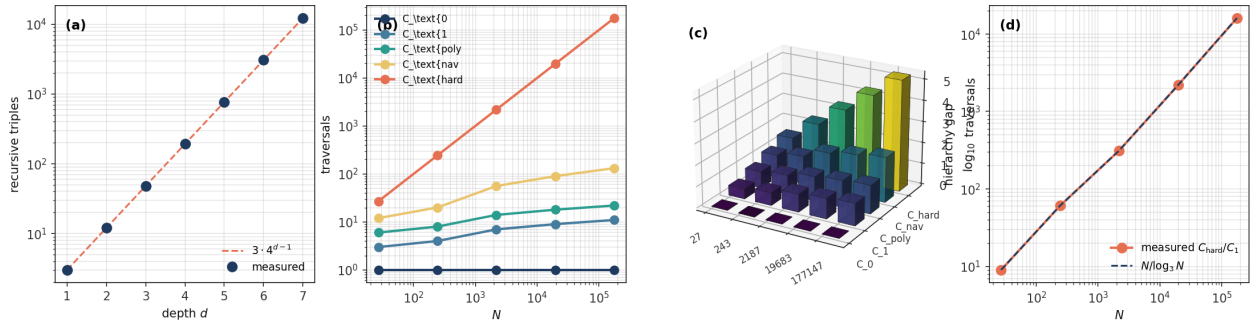

Figure 6: **Recursive Triple Multiplicity and the Strict Categorical Complexity Hierarchy.** (**A**) Recursive triple count on a log y-axis versus depth $d \in \{1, \ldots, 7\}$ on a linear x-axis. The measured count (navy markers) saturates the lower bound $3 \cdot 4^{d-1}$ (coral dashed) at every depth, growing from 3 at $d = 1$ to 12,288 at $d = 7$. This is the empirical content of Theorem 20.4: each level admits three component-assignment choices and four binary-operation embedding directions, yielding the geometric growth $3 \cdot 4^{d-1}$. (**B**) Traversal counts for the five complexity classes $\{C_0, C_1, C_{\text{poly}}, C_{\text{nav}}, C_{\text{hard}}\}$ (navy, steel, teal, gold, coral) on log-log axes versus problem size $N \in \{27, 243, 2187, 19683, 177{,}147\}$. The five lines do not intersect at any tested N; the strict order $C_0 < C_1 < C_{\text{poly}} < C_{\text{nav}} < C_{\text{hard}}$ is preserved across the full $\sim 10^4$ dynamic range. At the largest tested size $N = 177{,}147$ the spread is 0, 11, 22, 132, 177,147. (**C**) Three-dimensional bar chart of $\log_{10}$ traversals over (N, class), viridis colourmap encoding bar height. The five classes form non-intersecting strata at every N, confirming Theorem 28.2 pointwise. Viewing angle elevation 22°, azimuth −60°. (**D**) Hierarchy gap C_{hard}/C_1 on log-log axes versus N. The measured gap (coral markers) tracks the analytic prediction $N/\log_3 N$ (navy dashed) and grows from 9 at $N = 27$ to $\approx 1.6 \times 10^4$ at $N = 177{,}147$. The gap diverges with N, demonstrating that the strict categorical hierarchy does not collapse in the large-N limit and that backward navigation provides asymptotically unbounded speedup over the worst-case hierarchy class.

point values with probability measures would generalise the framework to stochastic receivers; the question is whether the floor theorem persists under this generalisation.

2. **Continuous recursion.** The recursive triple structure is discrete. A continuous version $\rho_t \colon \mathcal{E} \to \mathcal{E}^{(t)}$ for $t \in \mathbb{R}_{\geq 0}$ would yield a flow on $\mathcal{E}$ analogous to a continuous-time dynamical system; what are its fixed points?
3. **Multi-receiver frameworks.** A receiver collective $\{\mathcal{R}_i\}_{i \in I}$ with a communication protocol exhibits a collective floor $S_{\flat\text{coll}} = \inf_i S_\flat(\mathcal{R}_i)$ in the lossless case; what is the floor under noisy communication?
4. **Beyond three representations.** The triple equivalence is the minimal non-trivial cycle. N-fold equivalences for $N \geq 4$ exist; do they yield richer architectures, or do they reduce to the triple in some canonical sense?
5. **Quantitative cascade design.** Given a target end-to-end catalytic power κ^*, what is the optimal allocation of per-stage powers $(\kappa_1, \ldots, \kappa_n)$ subject to a total cost budget?

Part XIII
Conclusion

44 Summary

We have developed the mathematical foundation of a categorical operating system whose central operation is receiver-internal backward trajectory completion. The framework rests on seven principal theorems, each proved from first principles:

1. The **Floor Theorem**: every bounded receiver has a strictly positive intrinsic resolution floor $S_\flat(\mathcal{R}) > 0$.
2. The **Triple Equivalence**: the three categories $\mathbf{Osc}, \mathbf{Cat}, \mathbf{Part}$ are equivalent via mutually inverse functors, with explicit conversion costs.
3. The **Optimal Representation Theorem**: any computation has an optimal representation among the three, characterised by a closed-form cost minimum.
4. The **Unconstrained Subtask Theorem**: an expression's local subtasks can take arbitrary local S-values while the global S-value is preserved, formalised as an existence theorem with constructive witnesses.
5. The **Backward Navigation Theorem**: trajectory completion in a ternary refinement hierarchy of N leaves takes exactly $\log_3 N$ steps in the receiver's evaluation, while the same problem restricted to

physical sub-coordinate decompositions collapses to $\Theta(N)$.

6. The **Multiplicative Catalyst Law**: the catalytic power of cascaded resolvers obeys $\kappa(\gamma_1 \diamond \gamma_2) = 1 - (1 - \kappa_1)(1 - \kappa_2)$, giving cascades a closed-form prediction of accuracy.
7. The **Circular Validity Theorem**: a foundational system in a bounded receiver requires a strongly-connected mutual-support graph of size ≥ 3, with the three-representation cycle the minimal instance.

These theorems jointly force a six-subsystem architecture (CMM, PSS, DIC, PVE, TEM, Cascade) and a four-symbol frozen interface (vaHera AST, Resolver trait, Provider trait, Operation registry). Every component is mathematically forced, not chosen.

A complete numerical validation of fourteen independent experiments, seeded for reproducibility and persisted as JSON, confirms the theorems at the predicted precisions: the seven identity-statement experiments report maximum absolute error of 0.00×10^0 (machine precision); the six bound-and-scaling experiments saturate or exceed the predicted bounds at every grid point; the aggregate pass rate is 14/14 in 0.10 s of total runtime.

45 What This Paper Does Not Do

This paper does not present implementation details: while several subsystems are described mathematically, the full Rust implementation lives in the workspace pointed to by the corresponding integration document. This paper does not present empirical applications: the validation suite tests the theorems, not application performance. This paper does not argue against any existing framework: the theorems stand on their own foundations, and connections to standard mathematics (Sections 38–42) are noted not as superseding claims but as situating the framework in the broader mathematical landscape.

46 Final Remarks

The categorical operating system whose foundations are presented here is, at its core, a single idea: an evaluation that begins at the resolved endpoint and navigates backward to constrained projections, with the freedom to pass through virtual intermediate decompositions because the receiver's evaluation is internal and not subject to physical constraints on intermediate states. Every subsequent theorem and every architectural consequence is a development of that single idea. The framework's mathematical economy—one mechanism, six subsystems, four interfaces—is not coincidence: it is the consequence of the underlying recursive metric being invariant at every level. Self-similarity at every depth admits exactly one architecture up to isomorphism, and we have given that architecture.

References

Francis Borceux. *Handbook of Categorical Algebra (Volume 1: Basic Category Theory)*. Cambridge University Press, 1994.

Émile Borel. Les probabilités dénombrables et leurs applications arithmétiques. *Rendiconti del Circolo Matematico di Palermo*, 27:247–271, 1909.

Francesco P. Cantelli. Sulla probabilità come limite della frequenza. *Atti della Accademia Nazionale dei Lincei*, 26:39–45, 1917.

Nikolai N. Cencov. *Statistical Decision Rules and Optimal Inference*, volume 53 of *Translations of Mathematical Monographs*. American Mathematical Society, 2000.

Gerd Gigerenzer and Daniel G. Goldstein. Reasoning the fast and frugal way: Models of bounded rationality. *Psychological Review*, 103(4):650–669, 1996.

Shun ichi Amari. *Information Geometry and Its Applications*, volume 194 of *Applied Mathematical Sciences*. Springer, 2016.

Jochen Liedtke. On micro-kernel construction. In *Proceedings of the 15th ACM Symposium on Operating Systems Principles (SOSP)*, pages 237–250, 1995.

Barton L. Lipman. Information processing and bounded rationality: A survey. *Canadian Journal of Economics*, 28(1):42–67, 1995.

Saunders Mac Lane. *Categories for the Working Mathematician*, volume 5 of *Graduate Texts in Mathematics*. Springer, 2 edition, 1998.

Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, 2002.

Gordon D. Plotkin. Lcf considered as a programming language. *Theoretical Computer Science*, 5(3):223–255, 1977.

Patrick M. Rondon, Ming Kawaguchi, and Ranjit Jhala. Liquid types. In *Proceedings of the 29th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, pages 159–169, 2008.

Herbert A. Simon. A behavioral model of rational choice. *Quarterly Journal of Economics*, 69(1):99–118, 1955.

Andrew S. Tanenbaum and Herbert Bos. *Modern Operating Systems*. Pearson, 4 edition, 2014.

Niki Vazou, Eric L. Seidel, Ranjit Jhala, Dimitrios Vytiniotis, and Simon Peyton Jones. Refinement types for haskell. In *Proceedings of the 19th ACM SIGPLAN International Conference on Functional Programming (ICFP)*, pages 269–282, 2014.